

MP4_P1

November 16, 2020

1 Generative Adversarial Networks

For this part of the assignment you implement two different types of generative adversarial networks. We will train the networks on the Celeb A dataset which is a large set of celebrity face images.

```
[1]: import torch
from torch.utils.data import DataLoader
from torchvision import transforms
from torchvision.datasets import ImageFolder
from gan.utils import show_images
import numpy as np

%matplotlib inline
from matplotlib import pyplot as plt
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

%load_ext autoreload
%autoreload 2
```

```
[2]: from gan.train import train
```

```
[3]: device = torch.device("cuda:0" if torch.cuda.is_available() else "gpu")
```

2 GAN loss functions

In this assignment you will implement two different types of GAN cost functions. You will first implement the loss from the [original GAN paper](#). You will also implement the loss from [LS-GAN](#).

2.0.1 GAN loss

TODO: Implement the `discriminator_loss` and `generator_loss` functions in `gan/losses.py`.

The generator loss is given by:

$$\ell_G = -\mathbb{E}_{z \sim p(z)} [\log D(G(z))]$$

and the discriminator loss is:

$$\ell_D = -\mathbb{E}_{x \sim p_{\text{data}}} [\log D(x)] - \mathbb{E}_{z \sim p(z)} [\log (1 - D(G(z)))]$$

Note that these are negated from the equations presented earlier as we will be *minimizing* these losses.

HINTS: You should use the `torch.nn.functional.binary_cross_entropy_with_logits` function to compute the binary cross entropy loss since it is more numerically stable than using a softmax followed by BCE loss. The BCE loss is needed to compute the log probability of the true label given the logits output from the discriminator. Given a score $s \in \mathbb{R}$ and a label $y \in \{0, 1\}$, the binary cross entropy loss is

$$bce(s, y) = -y * \log(s) - (1 - y) * \log(1 - s)$$

Instead of computing the expectation of $\log D(G(z))$, $\log D(x)$ and $\log (1 - D(G(z)))$, we will be averaging over elements of the minibatch, so make sure to combine the loss by averaging instead of summing.

[4]: `from gan.losses import discriminator_loss, generator_loss`

2.0.2 Least Squares GAN loss

TODO: Implement the `ls_discriminator_loss` and `ls_generator_loss` functions in `gan/losses.py`.

We'll now look at [Least Squares GAN](#), a newer, more stable alternative to the original GAN loss function. For this part, all we have to do is change the loss function and retrain the model. We'll implement equation (9) in the paper, with the generator loss:

$$\ell_G = \frac{1}{2} \mathbb{E}_{z \sim p(z)} [(D(G(z)) - 1)^2]$$

and the discriminator loss:

$$\ell_D = \frac{1}{2} \mathbb{E}_{x \sim p_{\text{data}}} [(D(x) - 1)^2] + \frac{1}{2} \mathbb{E}_{z \sim p(z)} [(D(G(z)))^2]$$

HINTS: Instead of computing the expectation, we will be averaging over elements of the minibatch, so make sure to combine the loss by averaging instead of summing. When plugging in for $D(x)$ and $D(G(z))$ use the direct output from the discriminator (`scores_real` and `scores_fake`).

[5]: `from gan.losses import ls_discriminator_loss, ls_generator_loss`

3 GAN model architecture

TODO: Implement the Discriminator and Generator networks in `gan/models.py`.

We recommend the following architectures which are inspired by [DCGAN](#):

Discriminator:

- convolutional layer with in_channels=3, out_channels=128, kernel=4, stride=2
- convolutional layer with in_channels=128, out_channels=256, kernel=4, stride=2
- batch norm
- convolutional layer with in_channels=256, out_channels=512, kernel=4, stride=2
- batch norm
- convolutional layer with in_channels=512, out_channels=1024, kernel=4, stride=2
- batch norm
- convolutional layer with in_channels=1024, out_channels=1, kernel=4, stride=1

Instead of Relu we LeakyReLU throughout the discriminator (we use a negative slope value of 0.2).

The output of your discriminator should be a single value score corresponding to each input sample. See `torch.nn.LeakyReLU`.

Generator:

Note: In the generator, you will need to use transposed convolution (sometimes known as fractionally-strided convolution or deconvolution). This function is implemented in pytorch as `torch.nn.ConvTranspose2d`.

- transpose convolution with in_channels=NOISE_DIM, out_channels=1024, kernel=4, stride=1
- batch norm
- transpose convolution with in_channels=1024, out_channels=512, kernel=4, stride=2
- batch norm
- transpose convolution with in_channels=512, out_channels=256, kernel=4, stride=2
- batch norm
- transpose convolution with in_channels=256, out_channels=128, kernel=4, stride=2
- batch norm
- transpose convolution with in_channels=128, out_channels=3, kernel=4, stride=2

The output of the final layer of the generator network should have a tanh nonlinearity to output values between -1 and 1. The output should be a 3x64x64 tensor for each sample (equal dimensions to the images from the dataset).

```
[6]: import sys
sys.path.insert(0, "gan/")
from gan.models import Discriminator, Generator
```

4 Data loading: Celeb A Dataset

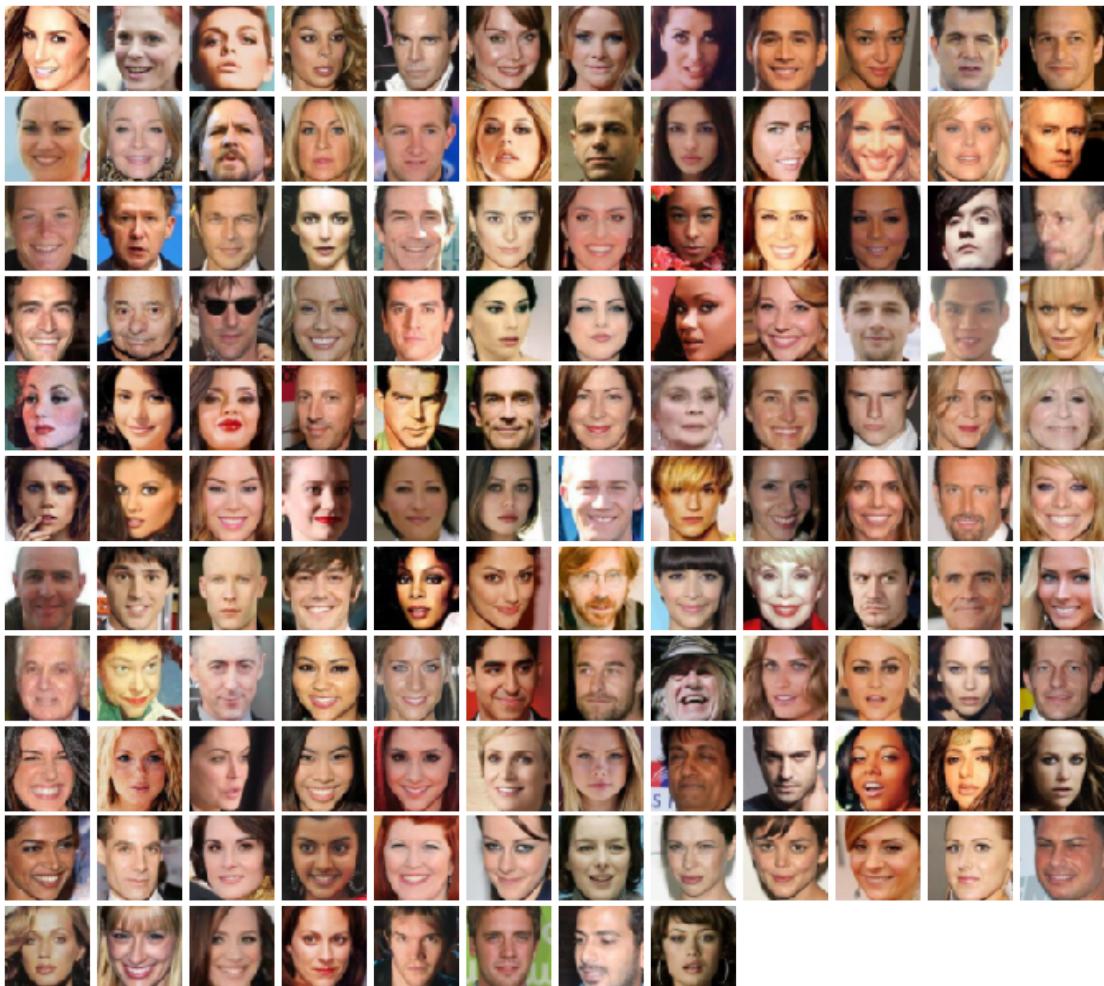
The CelebA images we provide have been filtered to obtain only images with clear faces and have been cropped and downsampled to 128x128 resolution.

```
[7]: batch_size = 128
scale_size = 64 # We resize the images to 64x64 for training
celeba_root = 'celeba_data'
```

```
[8]: celeba_train = ImageFolder(root=celeba_root, transform=transforms.Compose([
    transforms.Resize(scale_size),
    transforms.ToTensor(),
]))  
  
celeba_loader_train = DataLoader(celeba_train, batch_size=batch_size, u
→drop_last=True)
```

4.0.1 Visualize dataset

```
[9]: imgs = celeba_loader_train.__iter__().next()[0].numpy().squeeze()  
show_images(imgs, color=True)
```



5 Training

TODO: Fill in the training loop in gan/train.py.

```
[10]: NOISE_DIM = 100  
NUM_EPOCHS = 15  
learning_rate = 0.0002
```

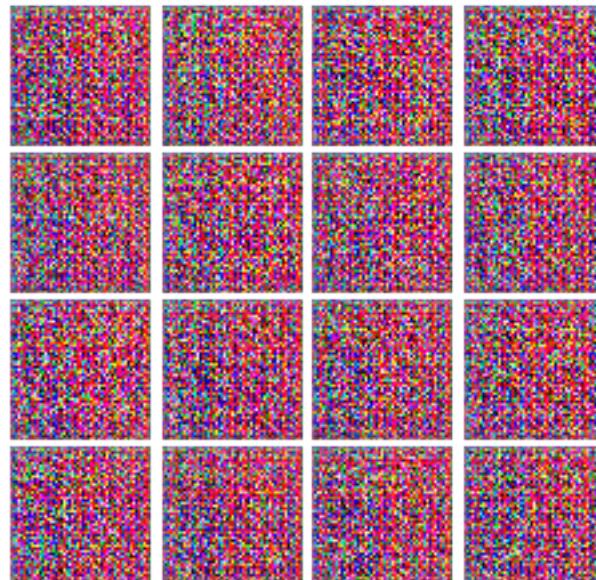
5.0.1 Train GAN

```
[11]: D = Discriminator().to(device)  
G = Generator(noise_dim=NOISE_DIM).to(device)
```

```
[12]: D_optimizer = torch.optim.Adam(D.parameters(), lr=learning_rate, betas = (0.5, 0.  
→999))  
G_optimizer = torch.optim.Adam(G.parameters(), lr=learning_rate, betas = (0.5, 0.  
→999))
```

```
[13]: # original gan  
train(D, G, D_optimizer, G_optimizer, discriminator_loss,  
      generator_loss, num_epochs=NUM_EPOCHS, show_every=150,  
      train_loader=celeba_loader_train, device=device)
```

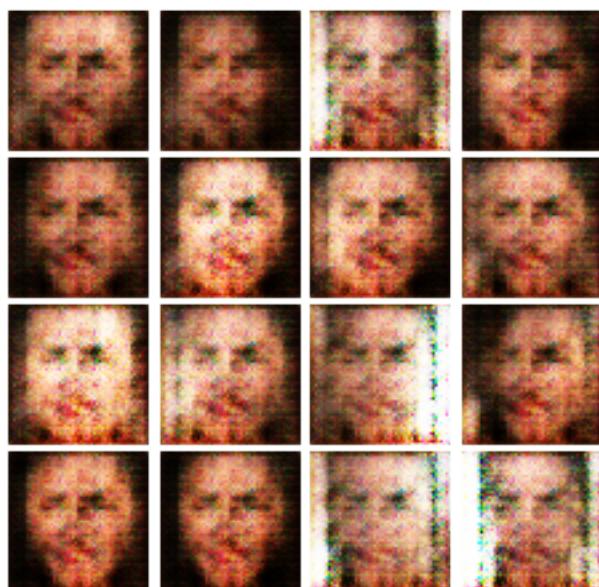
EPOCH: 1
Iter: 0, D: 1.272, G:11.46



Iter: 150, D: 0.8672, G:1.586



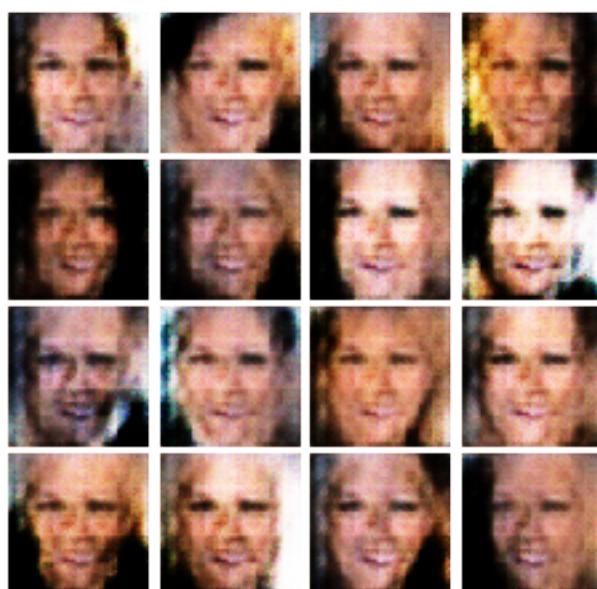
Iter: 300, D: 1.979, G:3.622



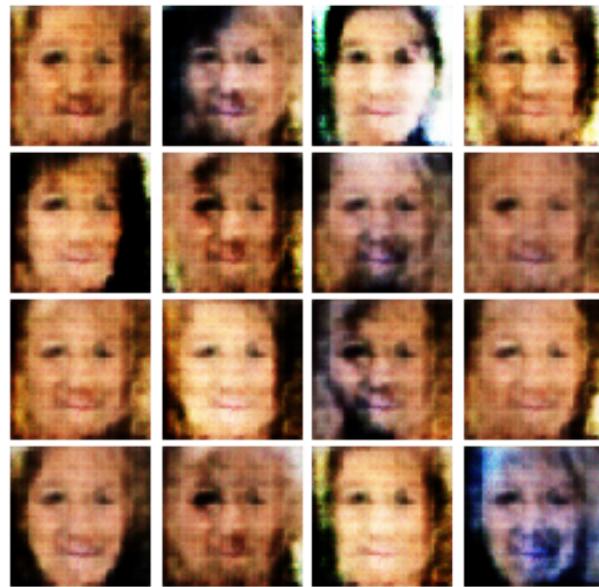
Iter: 450, D: 0.7691, G:2.214



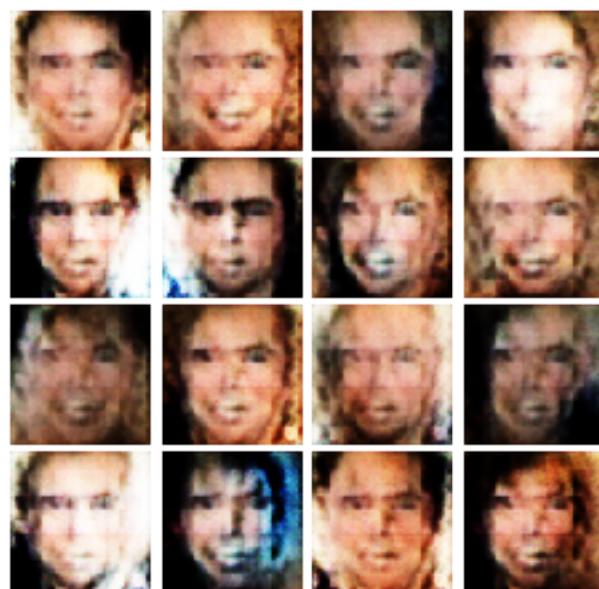
Iter: 600, D: 0.8871, G:4.267



Iter: 750, D: 0.5991, G:2.303

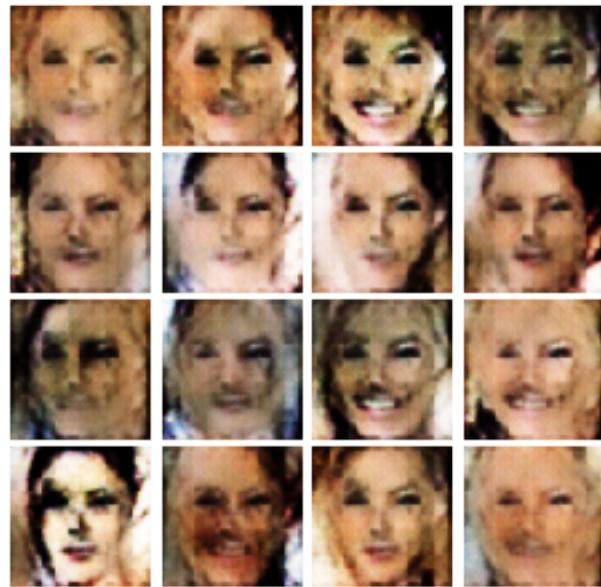


Iter: 900, D: 0.4315, G:3.691



EPOCH: 2

Iter: 1050, D: 0.4785, G:2.8



Iter: 1200, D: 0.7259, G:2.488



Iter: 1350, D: 0.4909, G:3.027



Iter: 1500, D: 0.7369, G:5.277



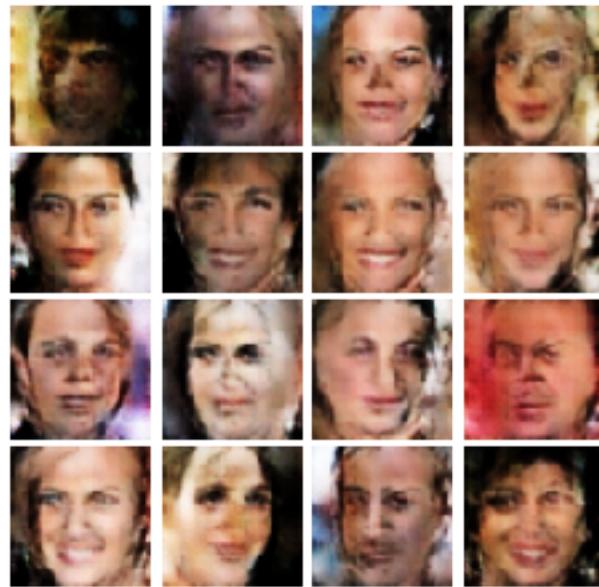
Iter: 1650, D: 0.7613, G:4.393



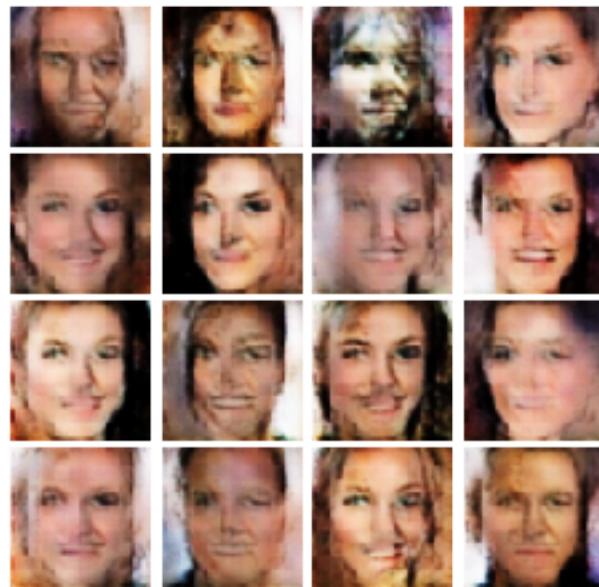
Iter: 1800, D: 0.2669, G:2.408



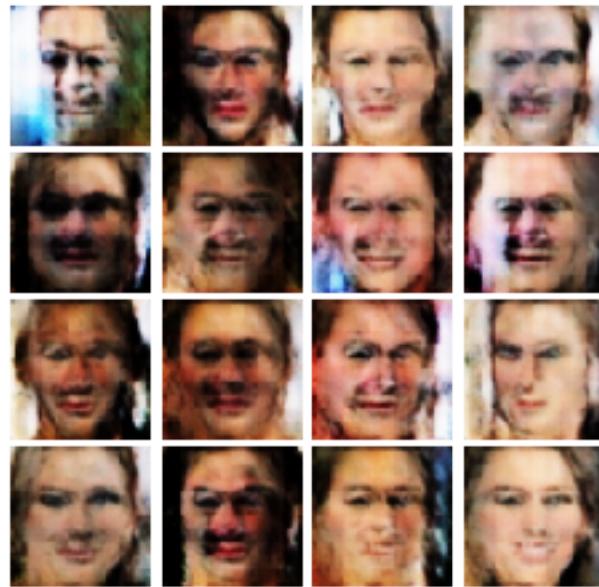
Iter: 1950, D: 0.1627, G:3.022



EPOCH: 3
Iter: 2100, D: 1.57, G:6.524



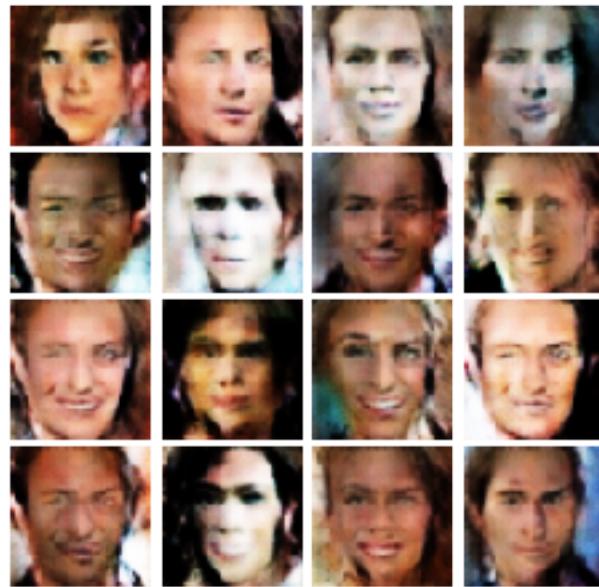
Iter: 2250, D: 0.9423, G:2.493



Iter: 2400, D: 0.1711, G:2.825



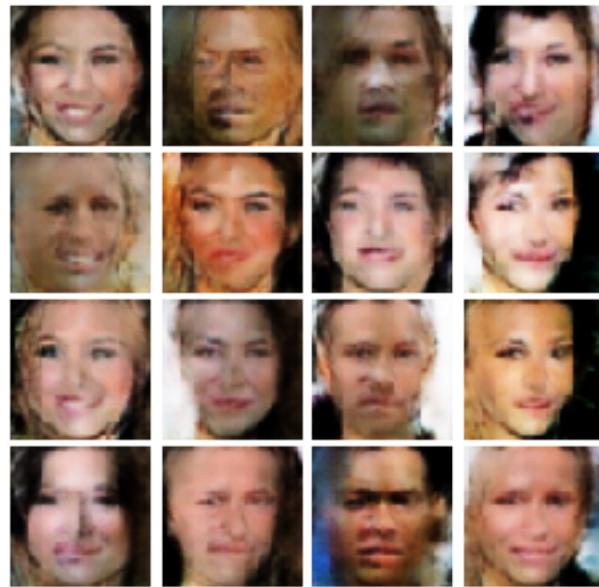
Iter: 2550, D: 0.138, G:2.724



Iter: 2700, D: 0.3795, G:3.612



Iter: 2850, D: 0.7126, G:5.387

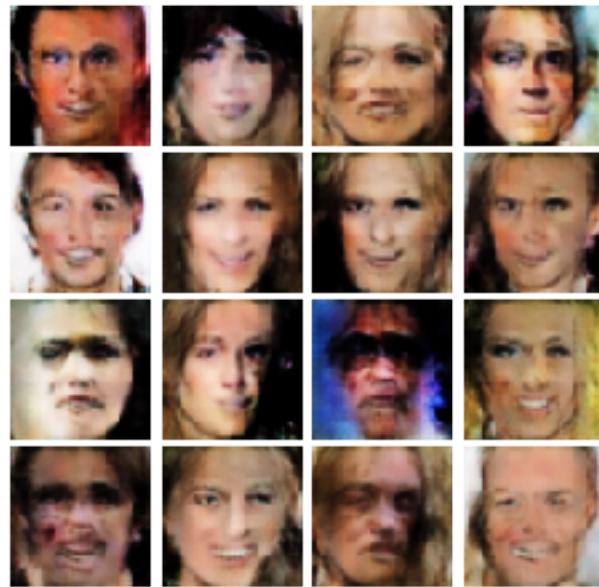


EPOCH: 4

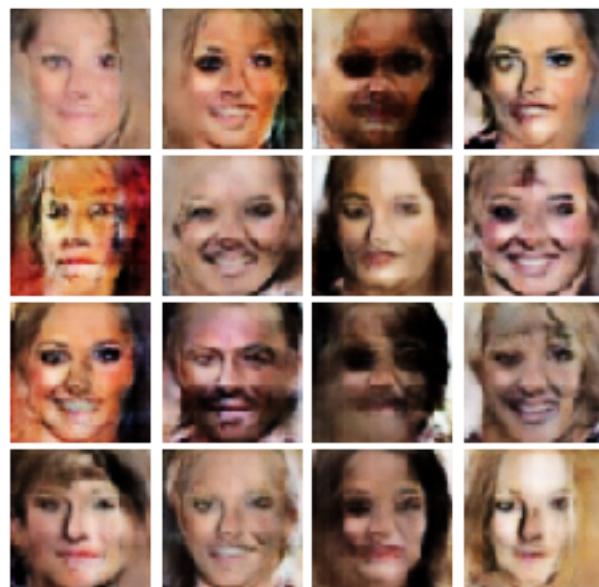
Iter: 3000, D: 0.4624, G: 4.246



Iter: 3150, D: 0.2361, G: 4.092



Iter: 3300, D: 0.04706, G:4.287



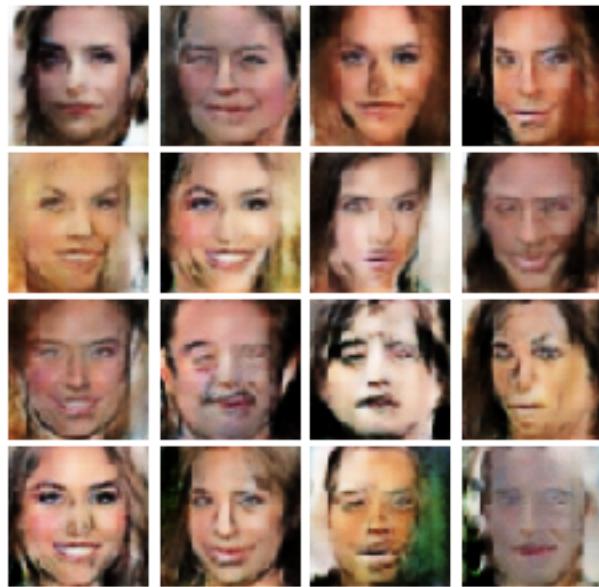
Iter: 3450, D: 0.04108, G:5.351



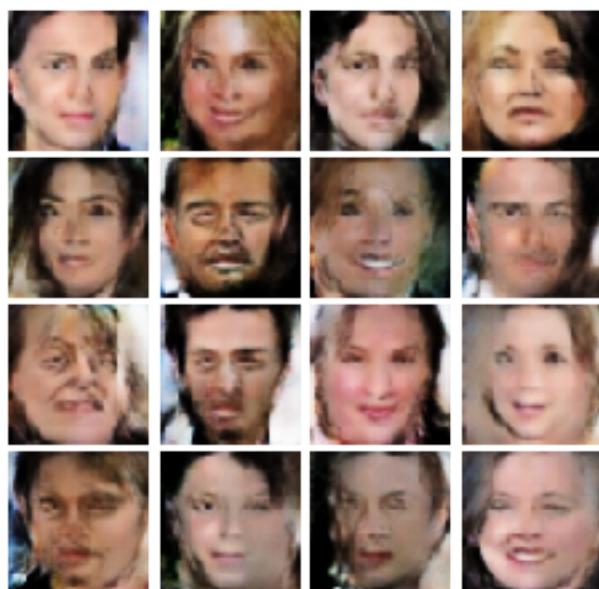
Iter: 3600, D: 0.4525, G: 4.946



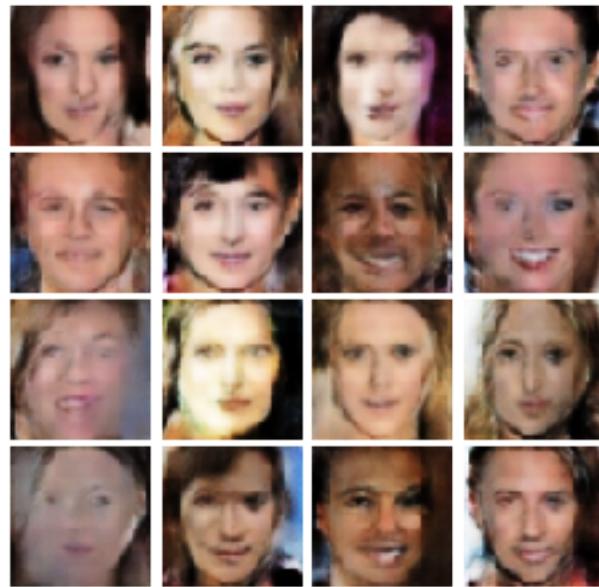
Iter: 3750, D: 0.7724, G: 6.248



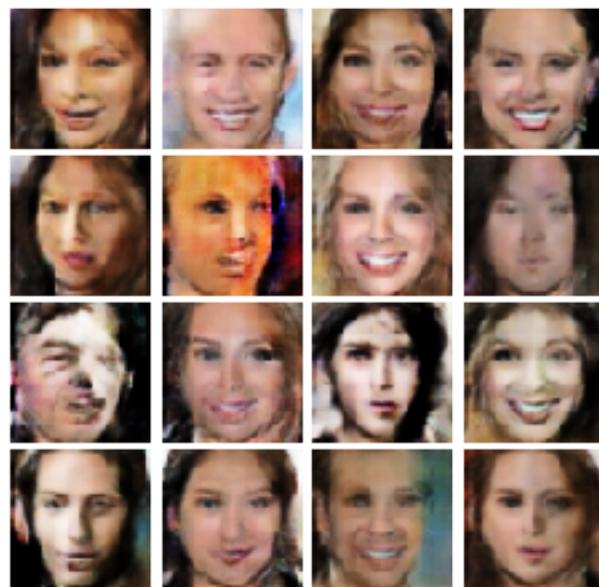
Iter: 3900, D: 0.1239, G:2.954



EPOCH: 5
Iter: 4050, D: 0.586, G:5.22



Iter: 4200, D: 0.05658, G:3.461



Iter: 4350, D: 0.05471, G:6.47



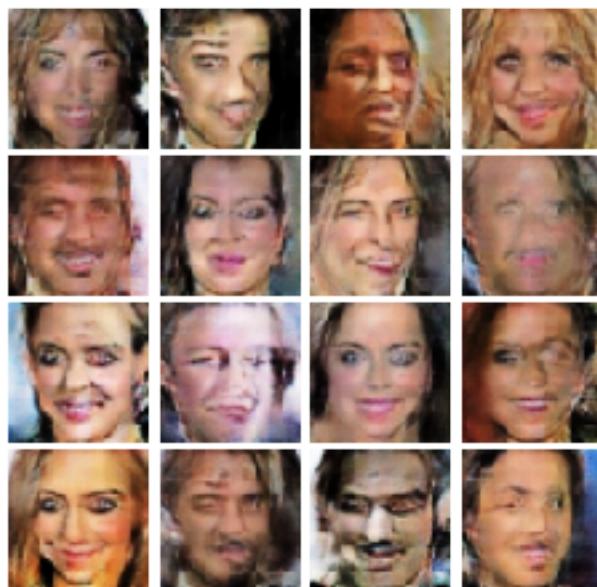
Iter: 4500, D: 0.33, G:3.028



Iter: 4650, D: 0.05662, G:3.678



Iter: 4800, D: 0.08579, G:1.966

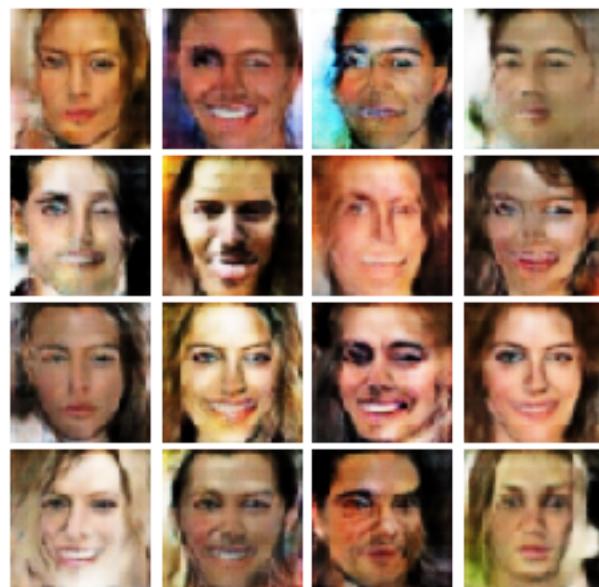


EPOCH: 6

Iter: 4950, D: 0.04087, G:4.593



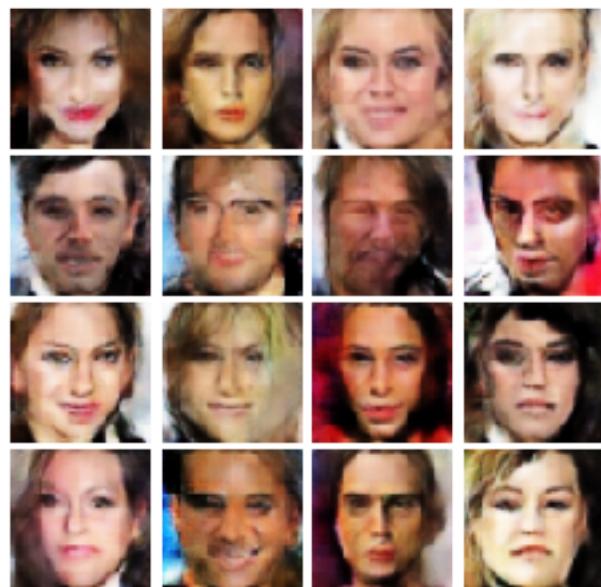
Iter: 5100, D: 0.1295, G:2.201



Iter: 5250, D: 0.0965, G:5.25



Iter: 5400, D: 0.07719, G:4.357



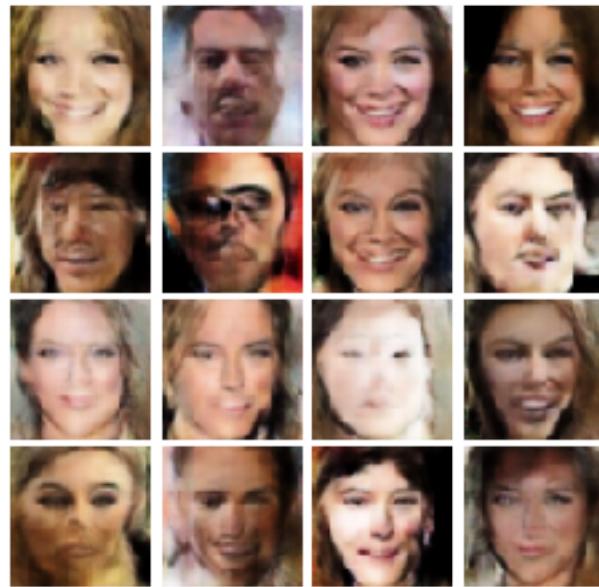
Iter: 5550, D: 0.1246, G:3.164



Iter: 5700, D: 0.4025, G: 5.581

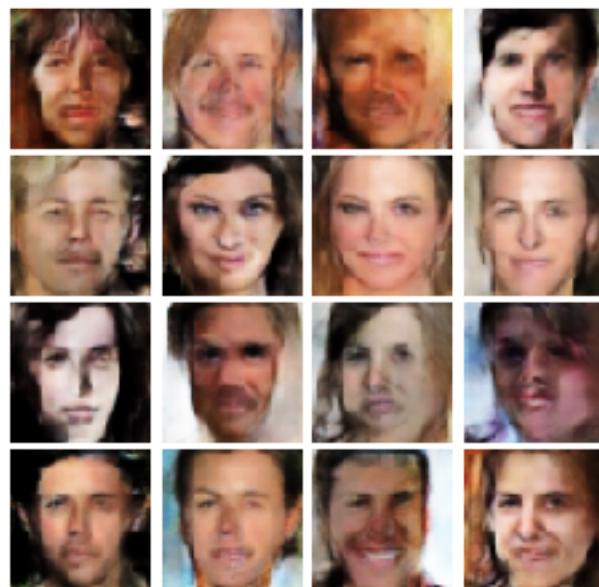


Iter: 5850, D: 0.1157, G: 2.967



EPOCH: 7

Iter: 6000, D: 0.1989, G:5.206



Iter: 6150, D: 2.444, G:4.699



Iter: 6300, D: 0.2108, G:5.43



Iter: 6450, D: 0.3089, G:3.936



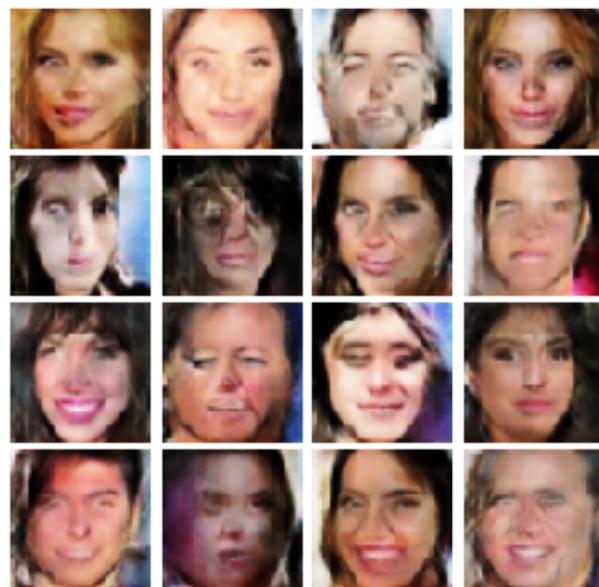
Iter: 6600, D: 0.5252, G:5.044



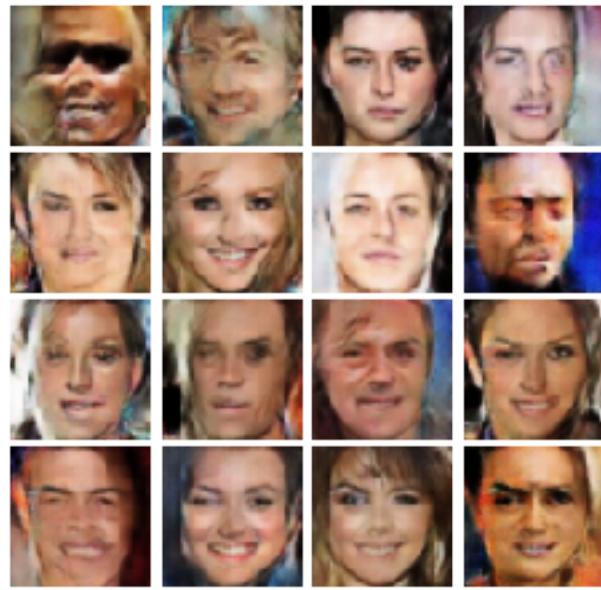
Iter: 6750, D: 0.04469, G:4.414



Iter: 6900, D: 0.6289, G:3.159



EPOCH: 8
Iter: 7050, D: 0.01218, G:5.025



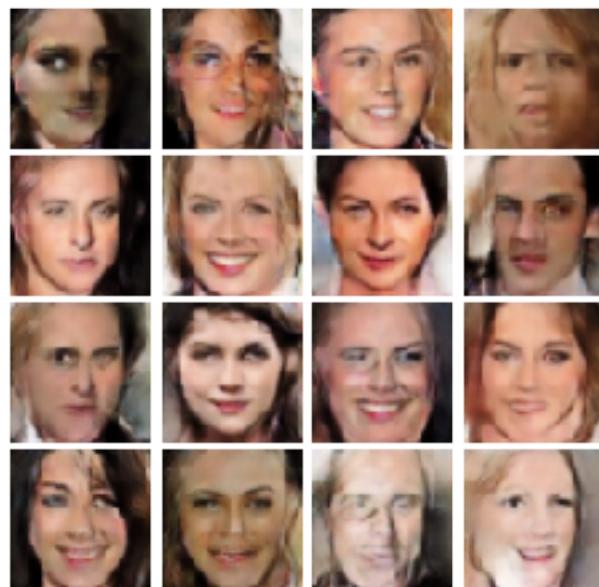
Iter: 7200, D: 0.2076, G: 4.304



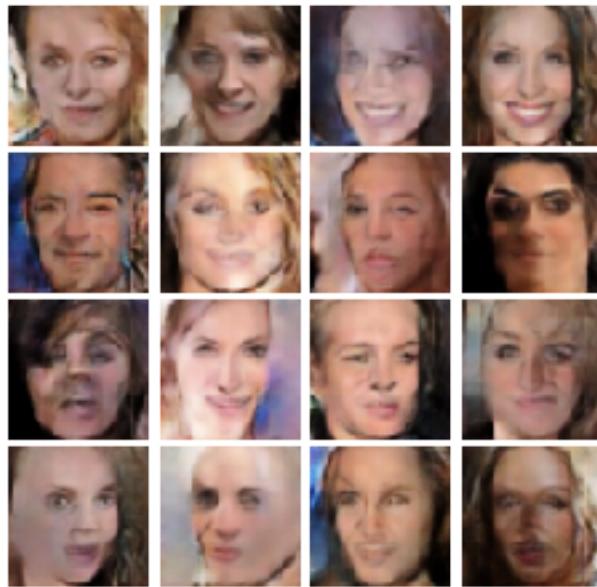
Iter: 7350, D: 0.1431, G: 3.173



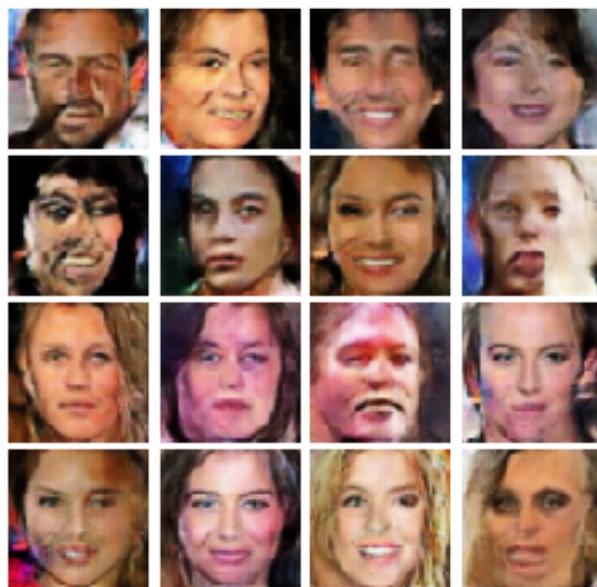
Iter: 7500, D: 0.07339, G:2.602



Iter: 7650, D: 0.3263, G:4.715

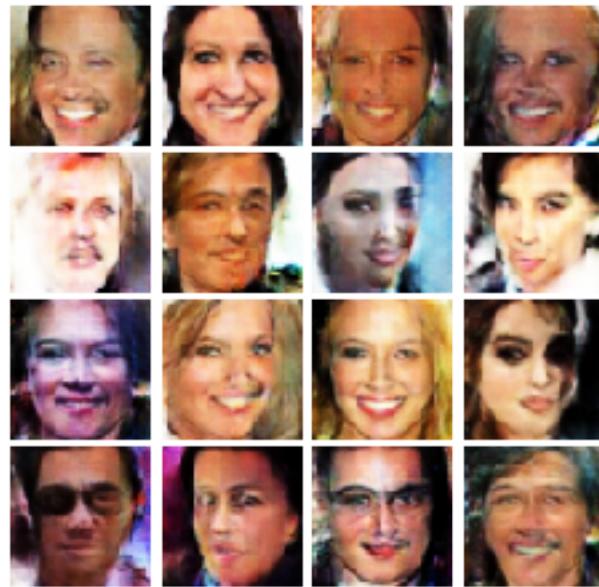


Iter: 7800, D: 0.1363, G: 5.405



EPOCH: 9

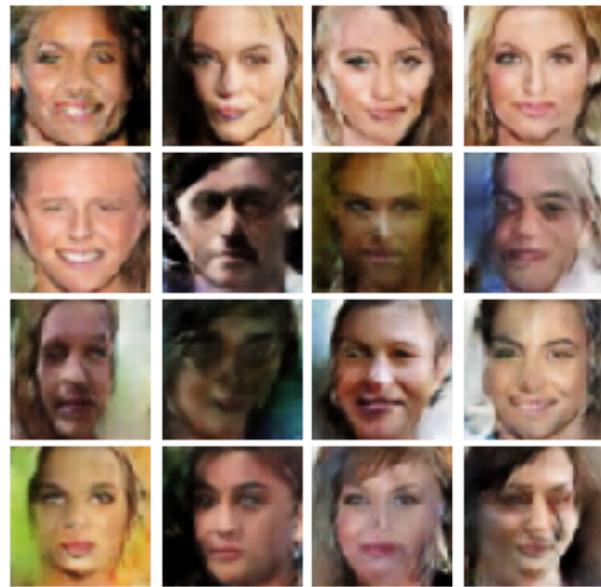
Iter: 7950, D: 0.09827, G: 4.293



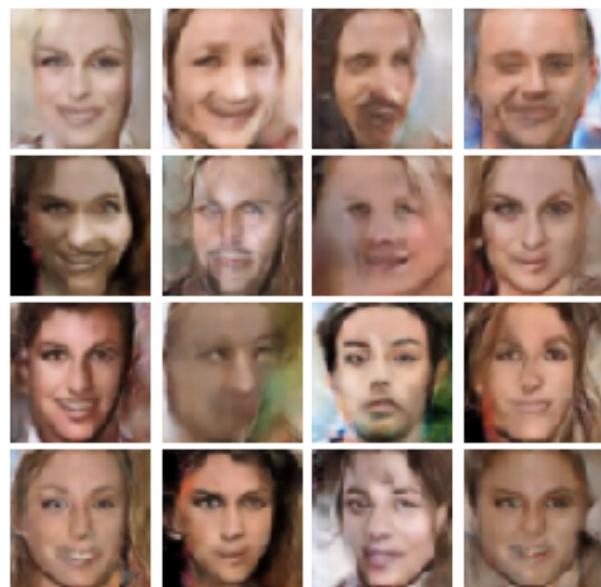
Iter: 8100, D: 0.09024, G:3.855



Iter: 8250, D: 0.1479, G:4.253



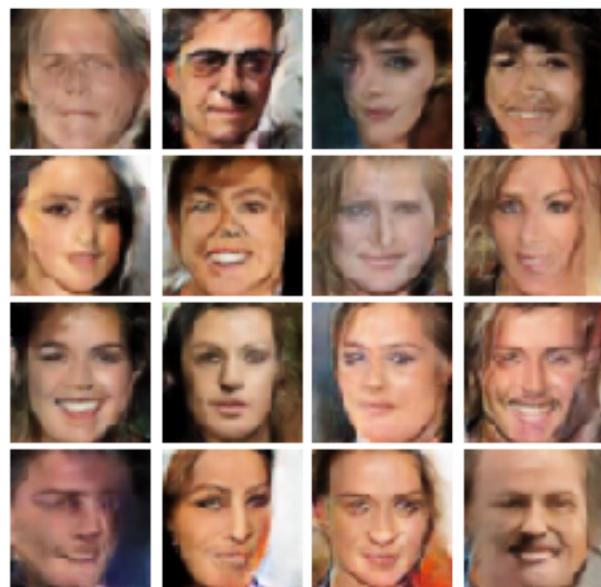
Iter: 8400, D: 0.3961, G:4.145



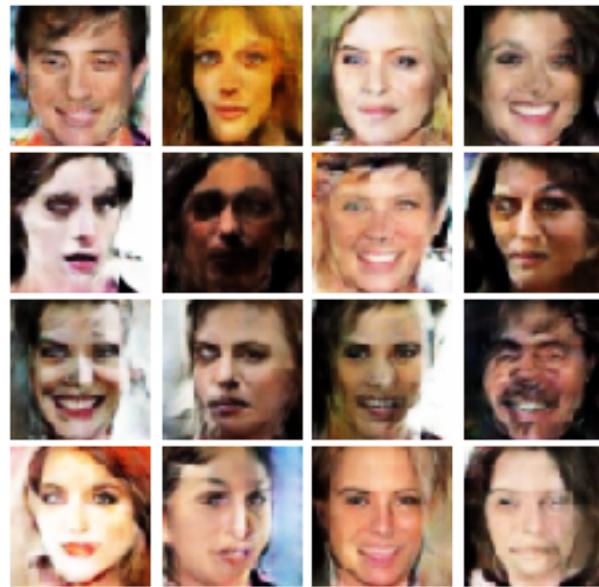
Iter: 8550, D: 0.3249, G:5.461



Iter: 8700, D: 0.7283, G:6.875

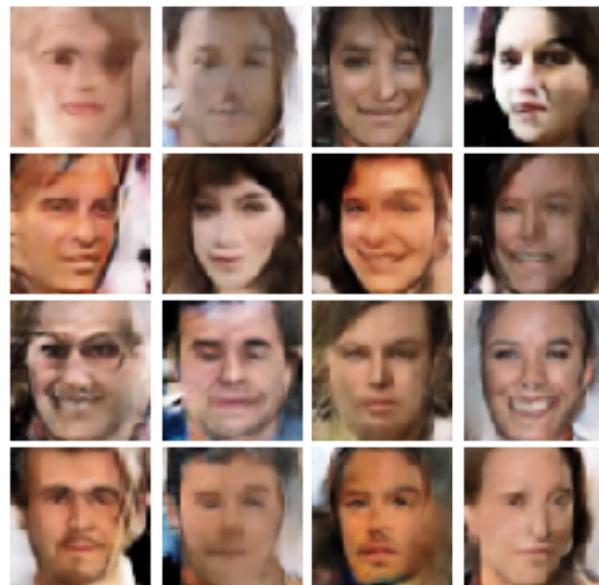


Iter: 8850, D: 0.1049, G:3.991

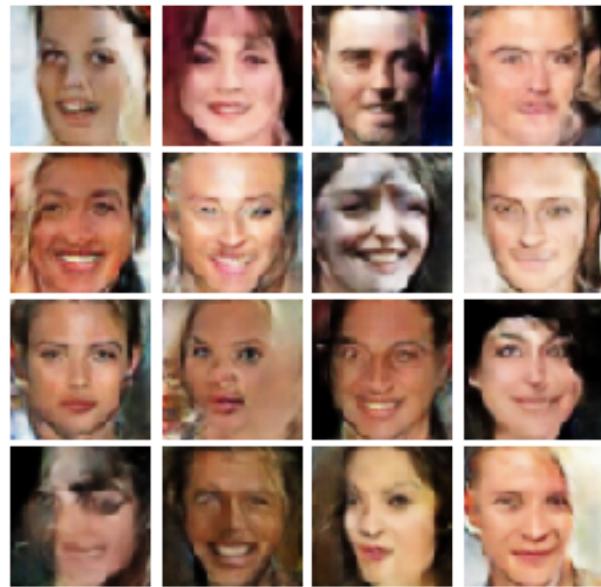


EPOCH: 10

Iter: 9000, D: 0.03767, G:4.114



Iter: 9150, D: 0.2483, G:4.04



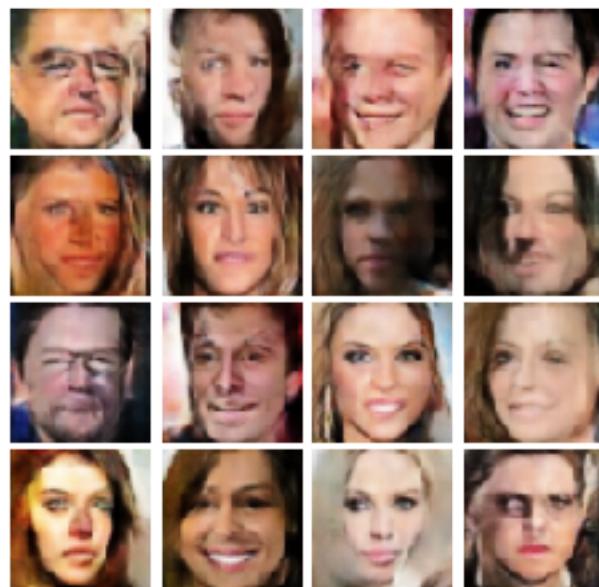
Iter: 9300, D: 0.1272, G:4.391



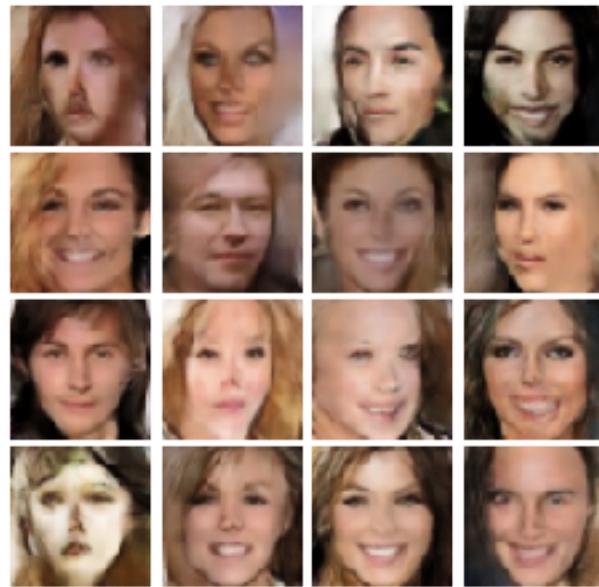
Iter: 9450, D: 0.2753, G:4.427



Iter: 9600, D: 0.1888, G:2.2

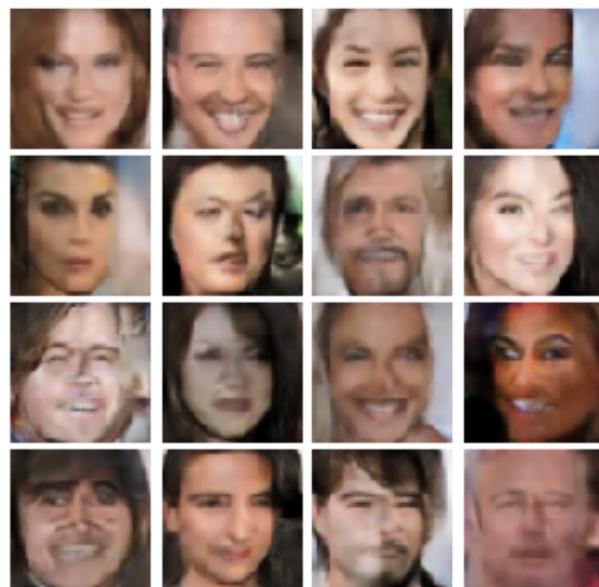


Iter: 9750, D: 0.3219, G:2.943

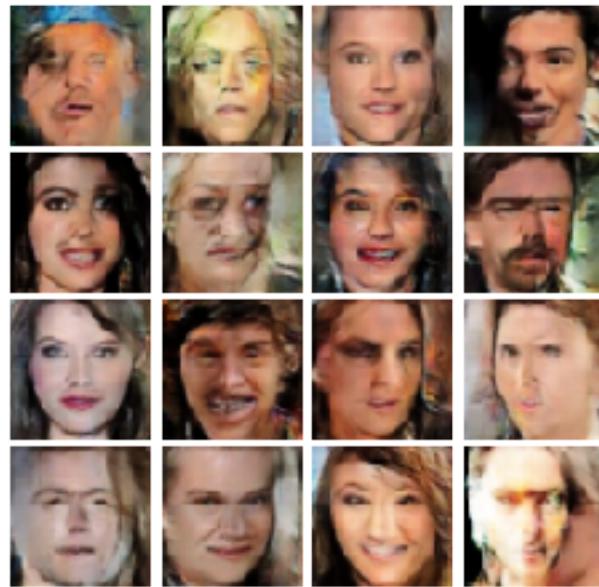


EPOCH: 11

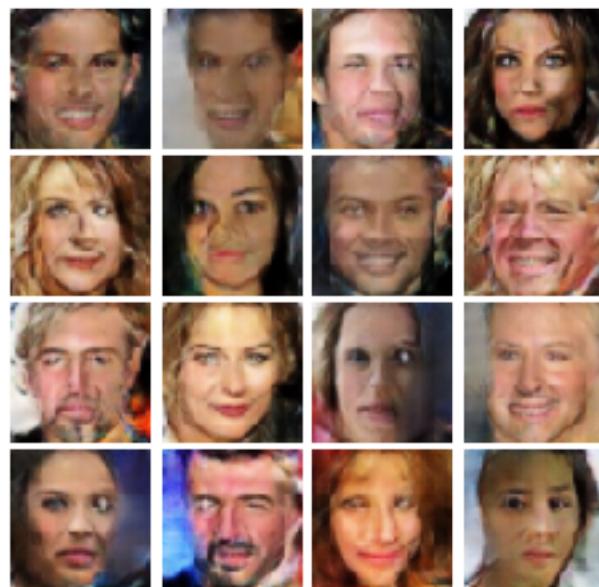
Iter: 9900, D: 0.1172, G: 4.978



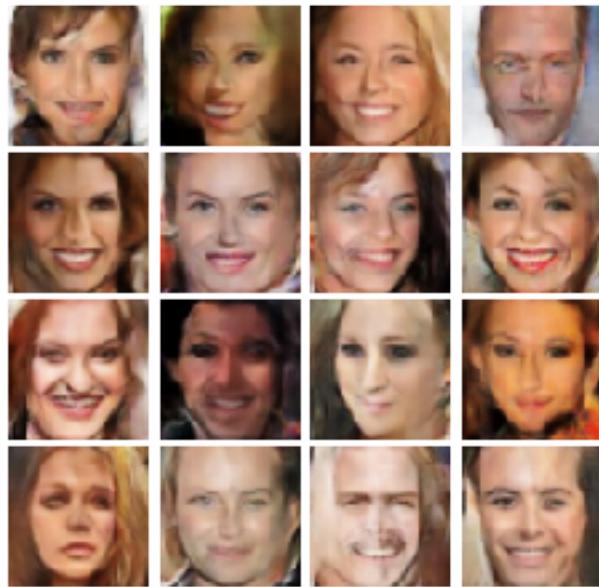
Iter: 10050, D: 0.2626, G: 3.738



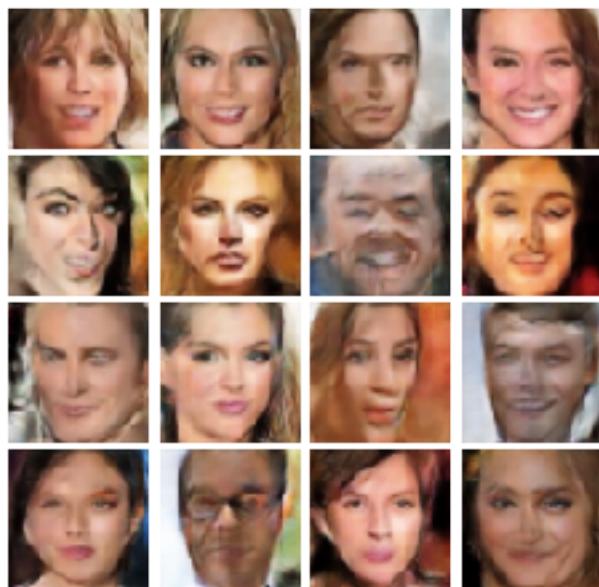
Iter: 10200, D: 0.1627, G:3.953



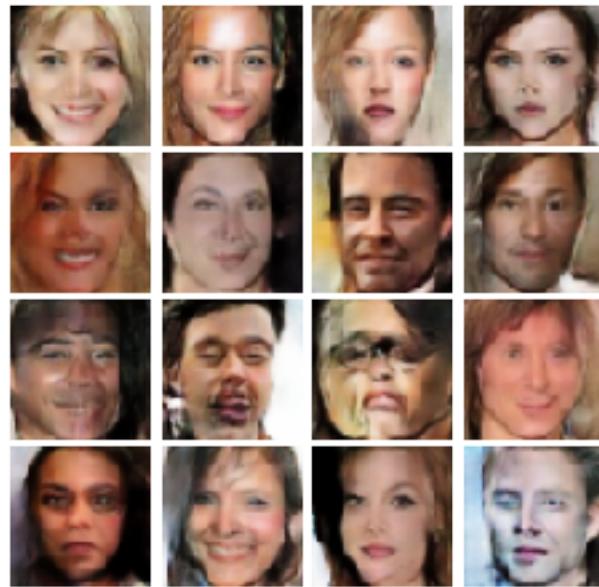
Iter: 10350, D: 5.619, G:5.265



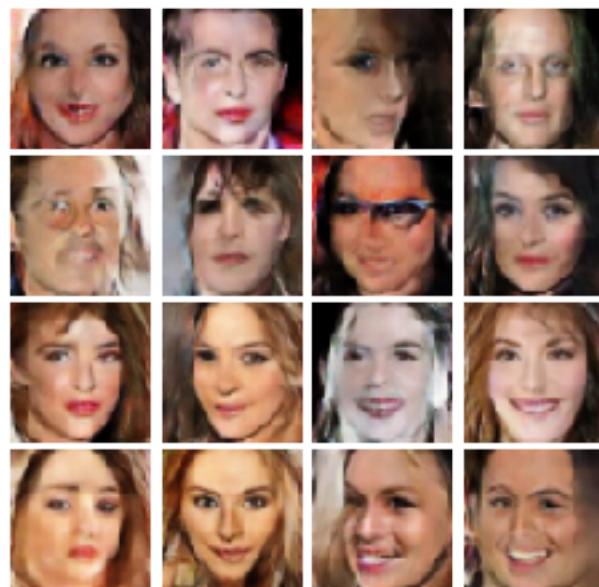
Iter: 10500, D: 0.1802, G:4.525



Iter: 10650, D: 0.1819, G:4.185

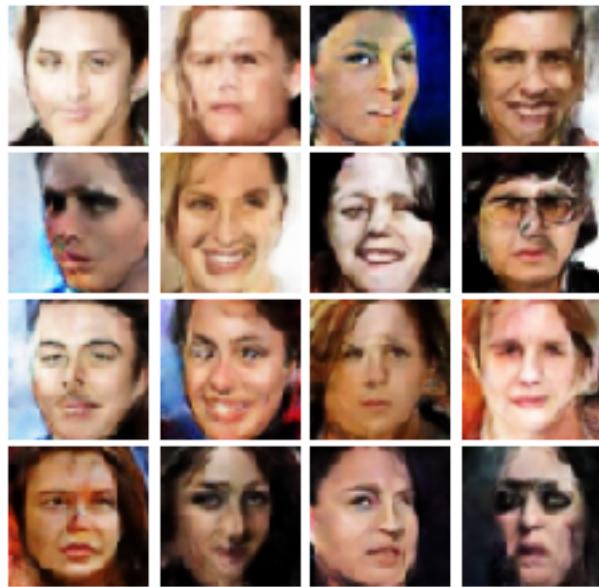


Iter: 10800, D: 0.04951, G:4.377



EPOCH: 12

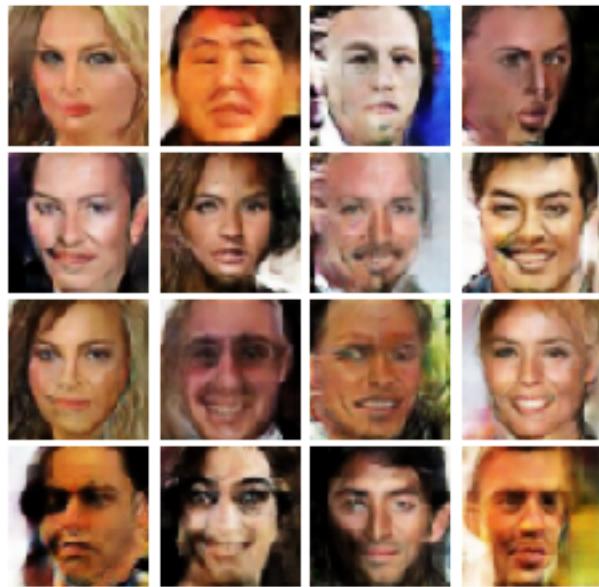
Iter: 10950, D: 0.05802, G:1.892



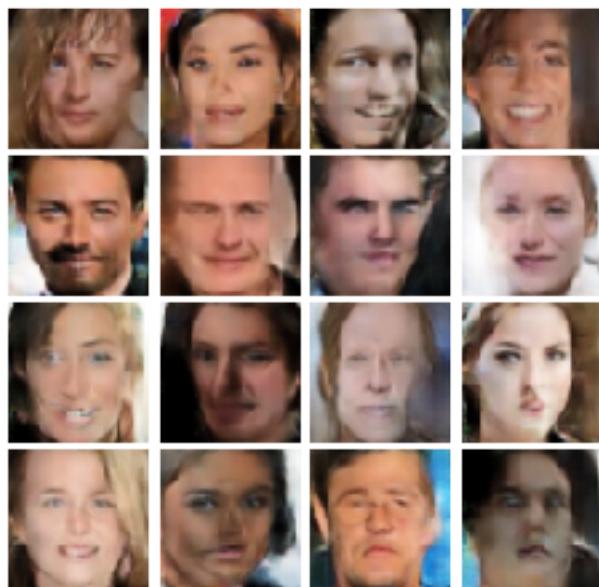
Iter: 11100, D: 0.2046, G:3.385



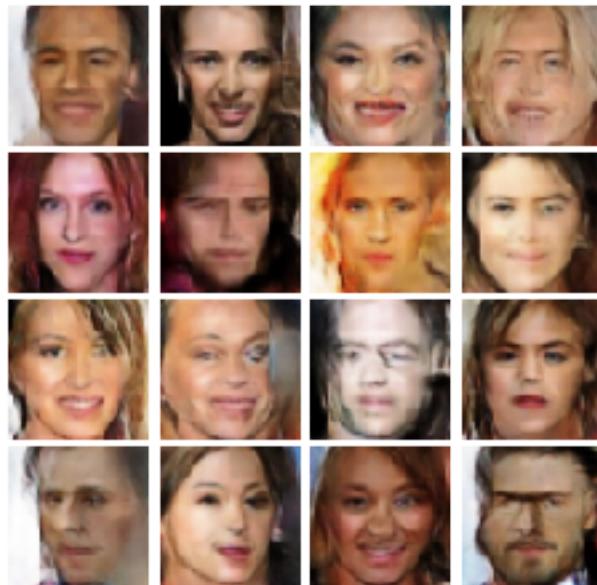
Iter: 11250, D: 0.1171, G:3.382



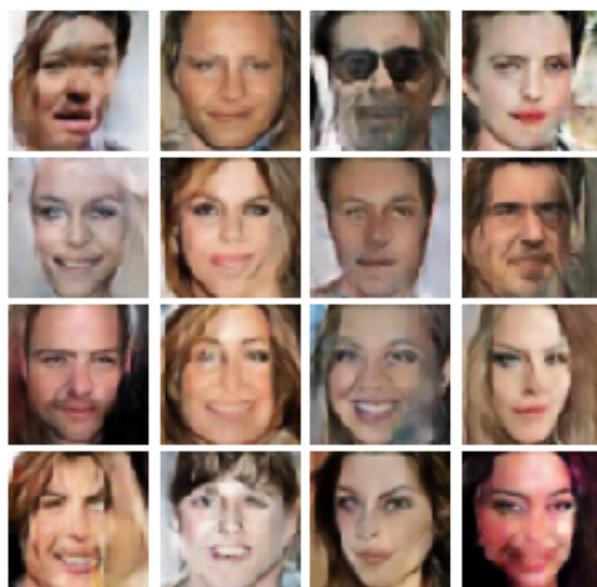
Iter: 11400, D: 0.3769, G: 5.405



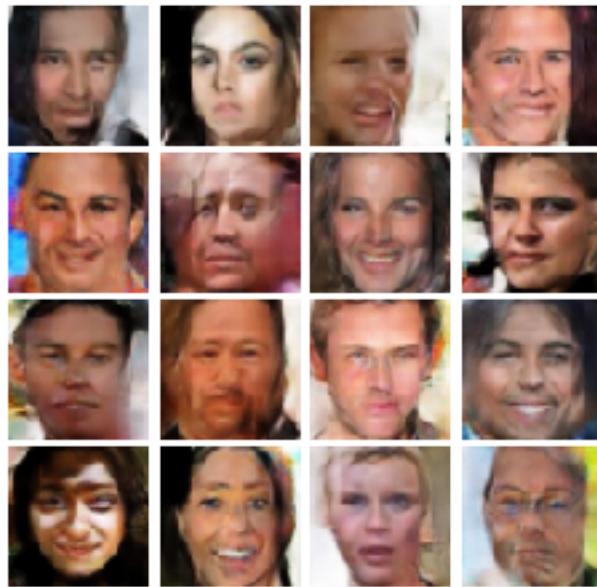
Iter: 11550, D: 0.09926, G: 4.508



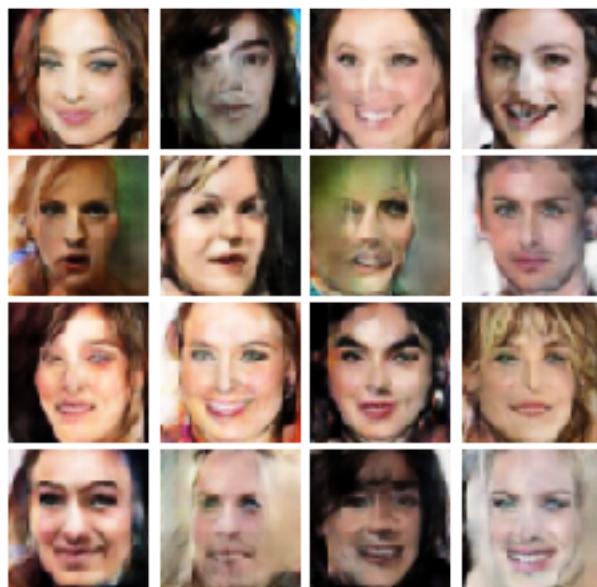
Iter: 11700, D: 0.1162, G:2.884



EPOCH: 13
Iter: 11850, D: 0.1642, G:5.728



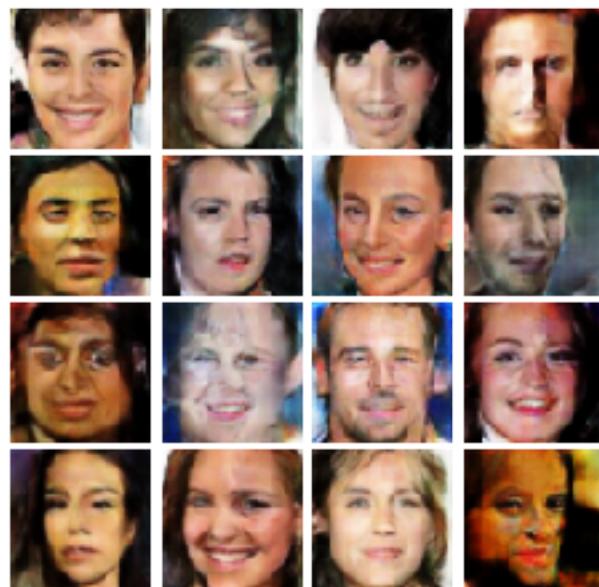
Iter: 12000, D: 0.1193, G:5.133



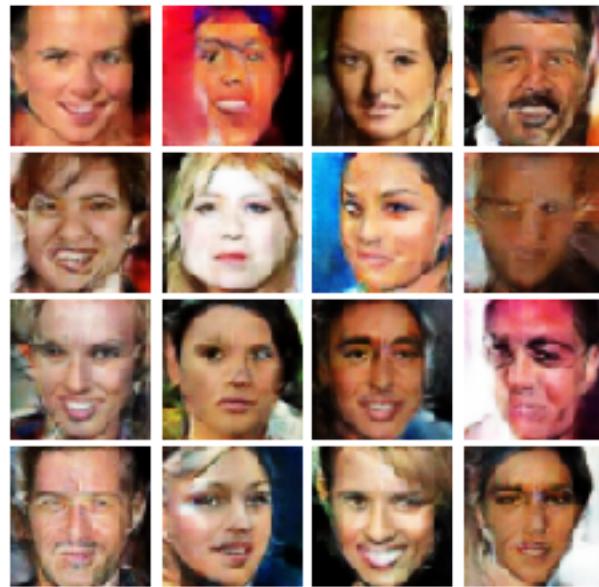
Iter: 12150, D: 0.2907, G:3.402



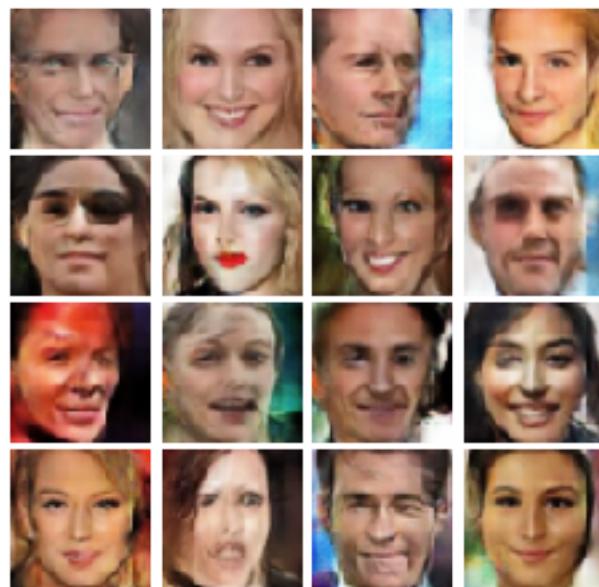
Iter: 12300, D: 0.1085, G:5.191



Iter: 12450, D: 0.1214, G:6.515



Iter: 12600, D: 0.7718, G: 4.171



Iter: 12750, D: 1.016, G: 9.094

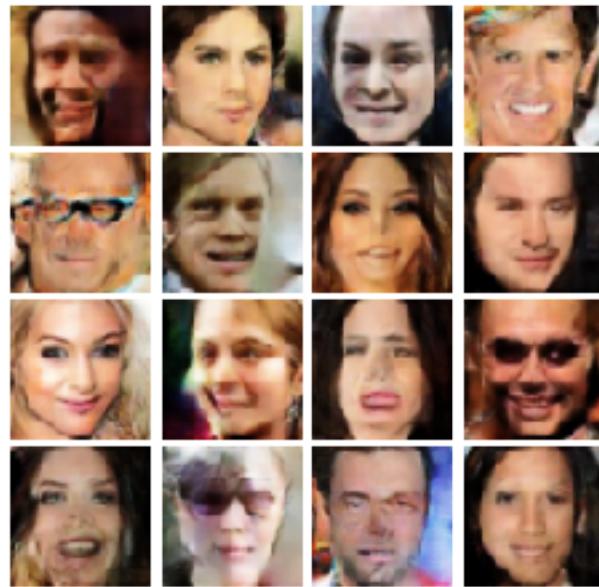


EPOCH: 14

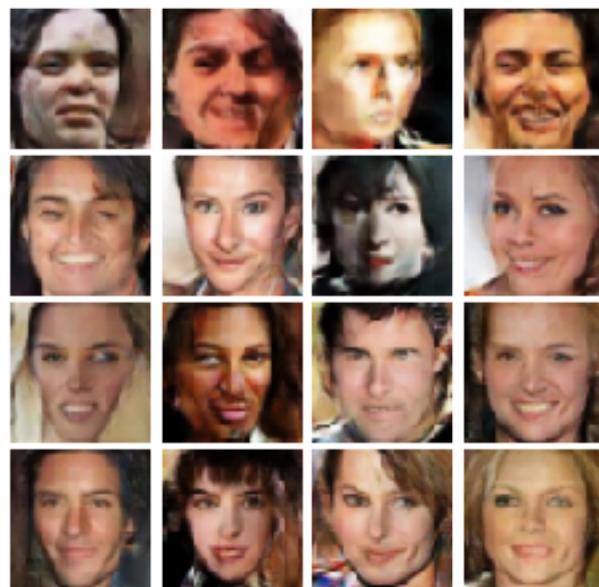
Iter: 12900, D: 0.05554, G: 4.271



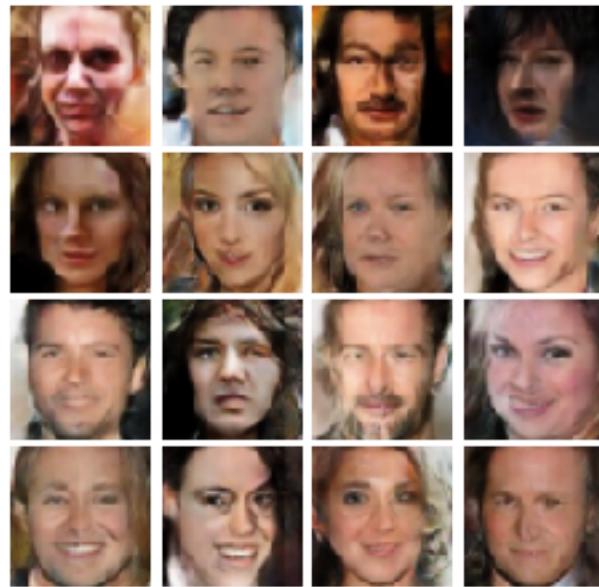
Iter: 13050, D: 1.605, G: 1.306



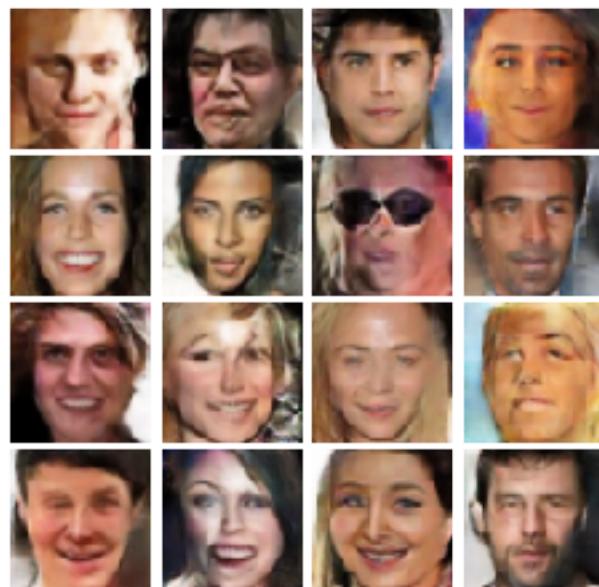
Iter: 13200, D: 0.2565, G:4.921



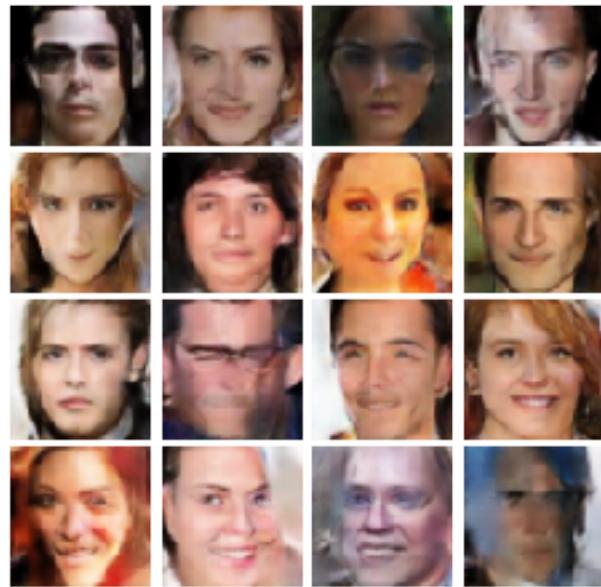
Iter: 13350, D: 0.9365, G:3.736



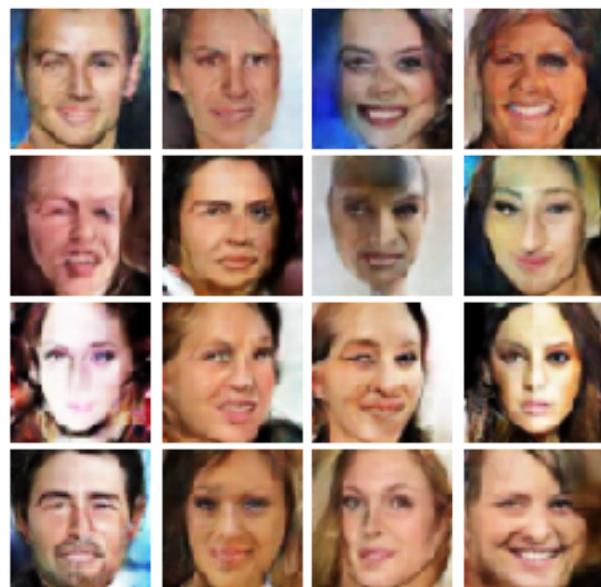
Iter: 13500, D: 0.3221, G:2.356



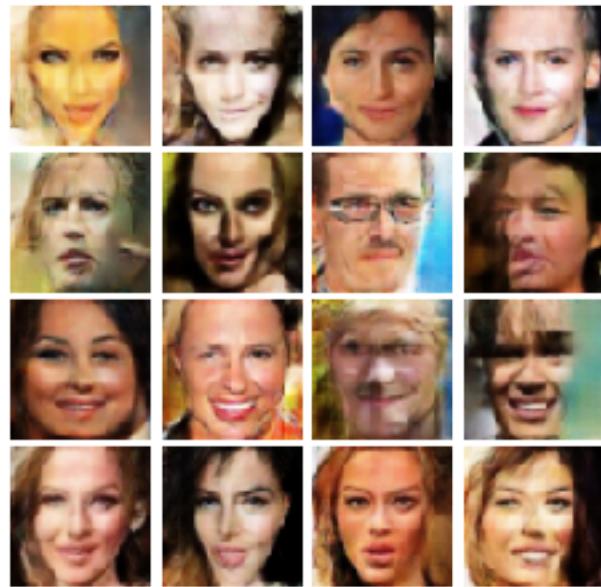
Iter: 13650, D: 0.1428, G:3.159



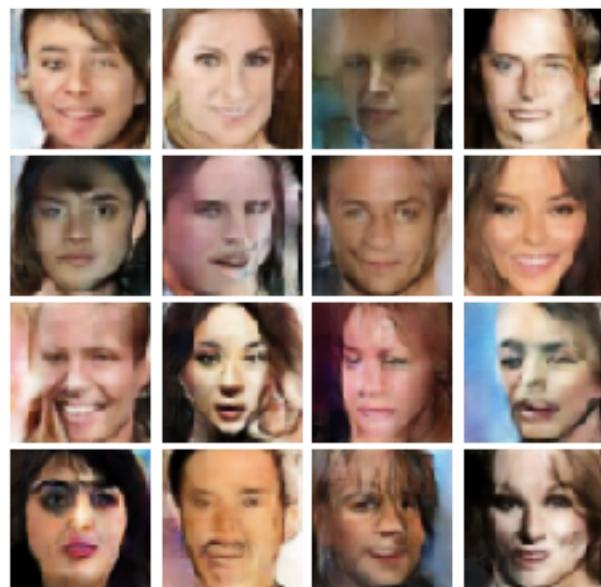
Iter: 13800, D: 0.2145, G: 5.252



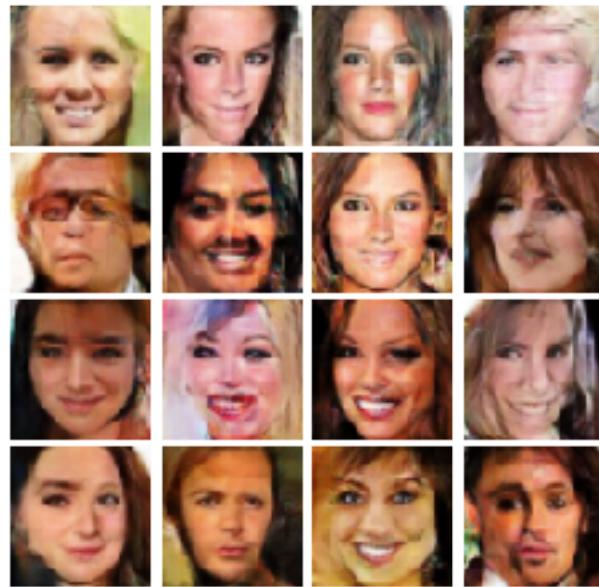
EPOCH: 15
Iter: 13950, D: 0.353, G: 3.559



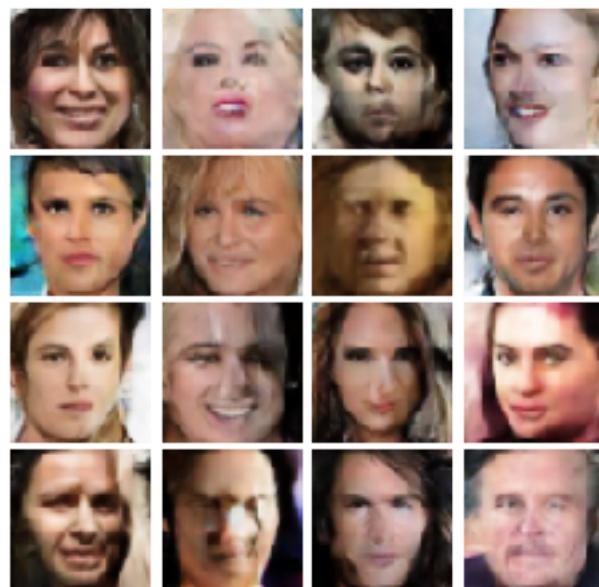
Iter: 14100, D: 0.2326, G:5.819



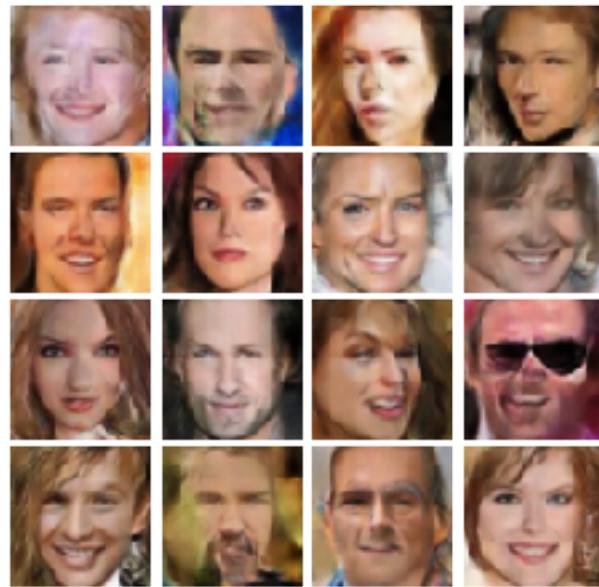
Iter: 14250, D: 0.1462, G:4.458



Iter: 14400, D: 1.052, G:1.516



Iter: 14550, D: 0.2075, G:3.163



Iter: 14700, D: 0.3034, G: 3.74



5.0.2 Train LS-GAN

```
[14]: D = Discriminator().to(device)
```

```
G = Generator(noise_dim=NOISE_DIM).to(device)
```

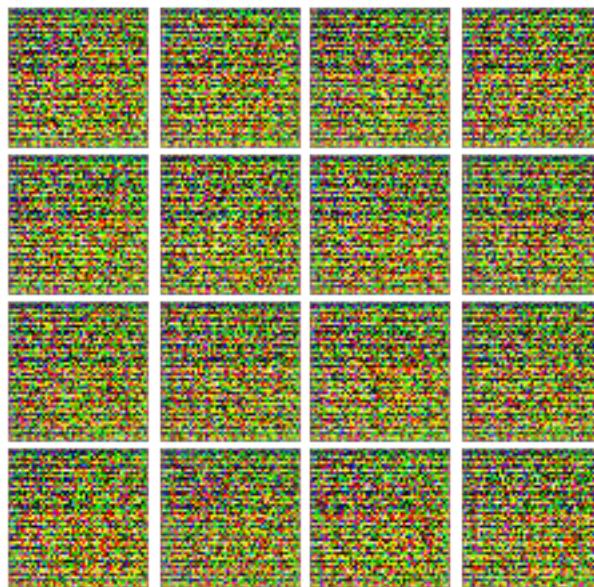
```
[15]: D_optimizer = torch.optim.Adam(D.parameters(), lr=learning_rate, betas = (0.5, 0.  
→999))
```

```
G_optimizer = torch.optim.Adam(G.parameters(), lr=learning_rate, betas = (0.5, 0.  
→999))
```

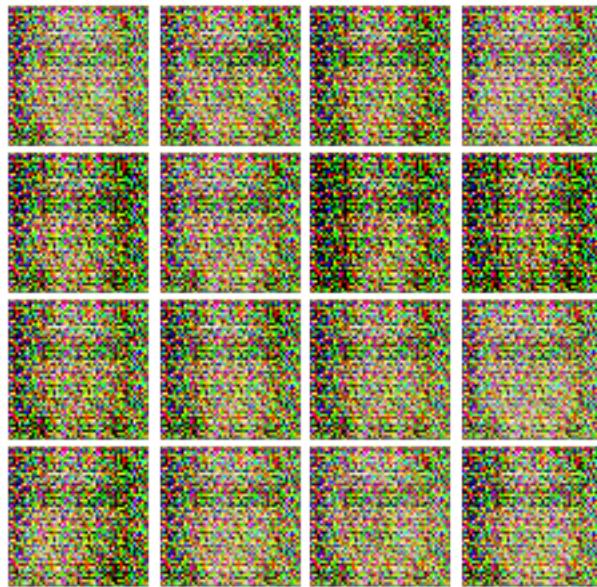
```
[16]: # ls-gan  
train(D, G, D_optimizer, G_optimizer, ls_discriminator_loss,  
      ls_generator_loss, num_epochs=NUM_EPOCHS, show_every=200,  
      train_loader=celeba_loader_train, device=device)
```

EPOCH: 1

Iter: 0, D: 1.126, G:21.14



Iter: 200, D: 0.08874, G:0.2953



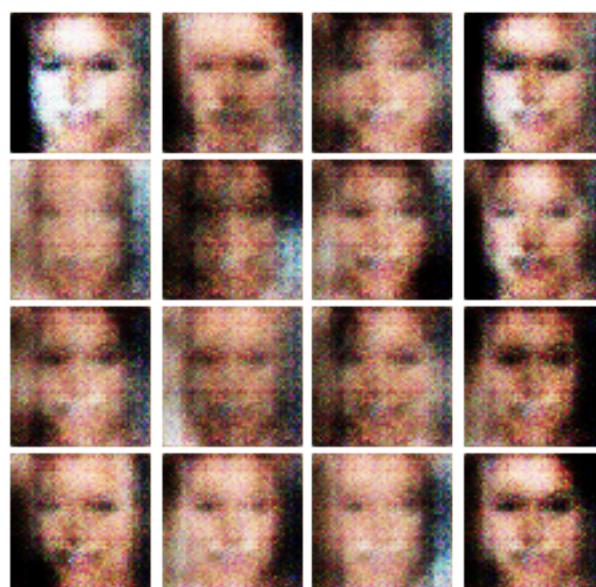
Iter: 400, D: 0.1626, G:0.128



Iter: 600, D: 0.09781, G:0.7947



Iter: 800, D: 0.1686, G: 0.7874



EPOCH: 2

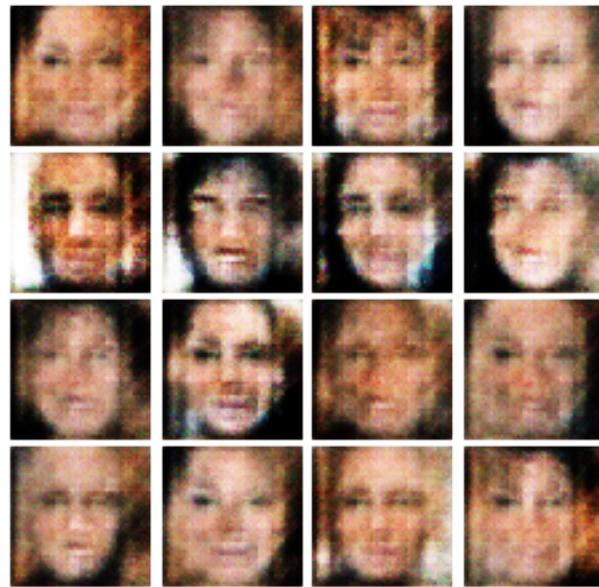
Iter: 1000, D: 0.2602, G: 0.6705



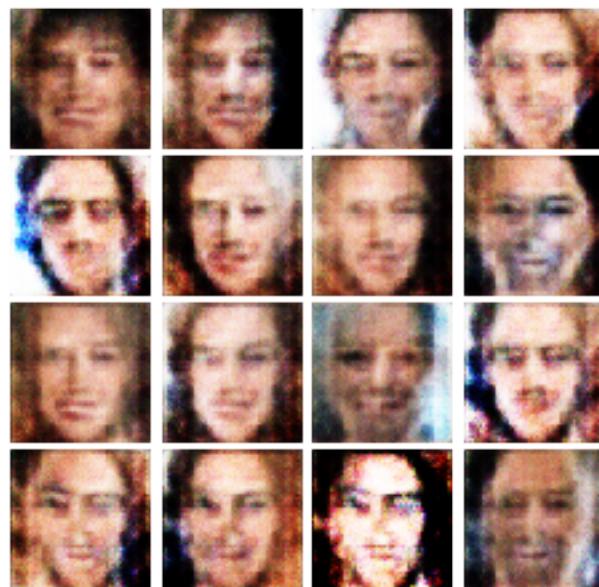
Iter: 1200, D: 0.2366, G: 0.1351



Iter: 1400, D: 0.1179, G: 0.2623



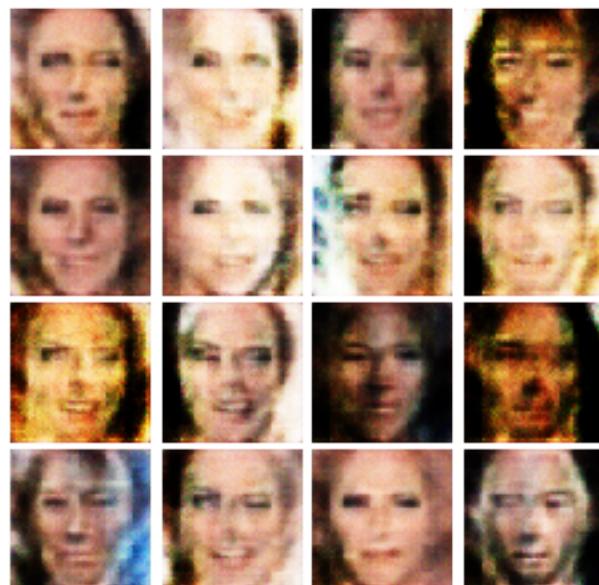
Iter: 1600, D: 0.1137, G: 0.1889



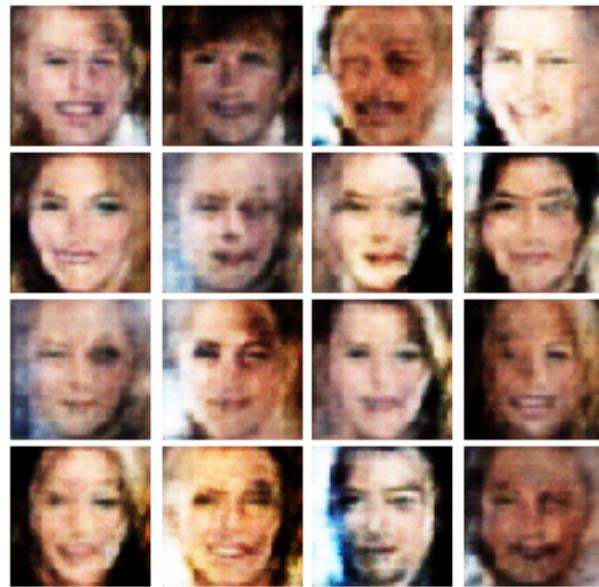
Iter: 1800, D: 0.2541, G: 0.2395



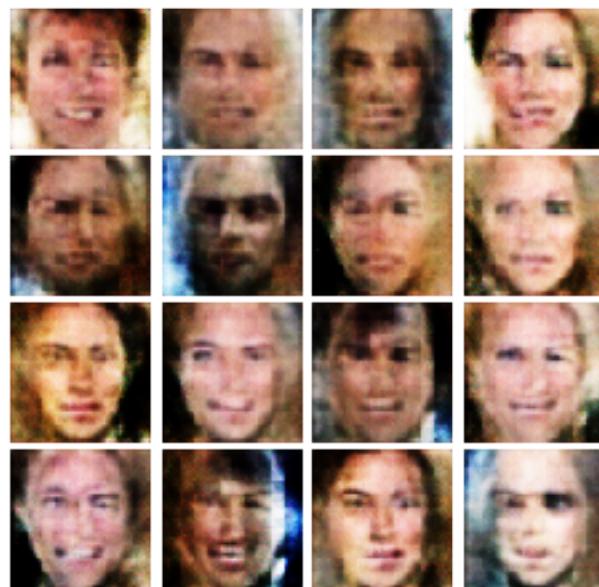
EPOCH: 3
Iter: 2000, D: 0.2508, G:0.3151



Iter: 2200, D: 0.1136, G:0.4612



Iter: 2400, D: 0.4136, G: 0.8757



Iter: 2600, D: 0.1284, G: 0.2979



Iter: 2800, D: 0.4212, G: 0.4918



EPOCH: 4

Iter: 3000, D: 0.4043, G: 0.08354



Iter: 3200, D: 0.1756, G:0.33



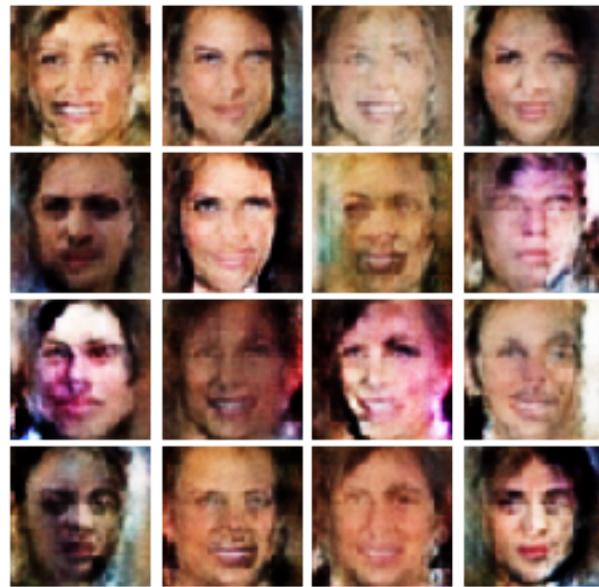
Iter: 3400, D: 0.356, G:1.239



Iter: 3600, D: 0.2411, G: 0.2179



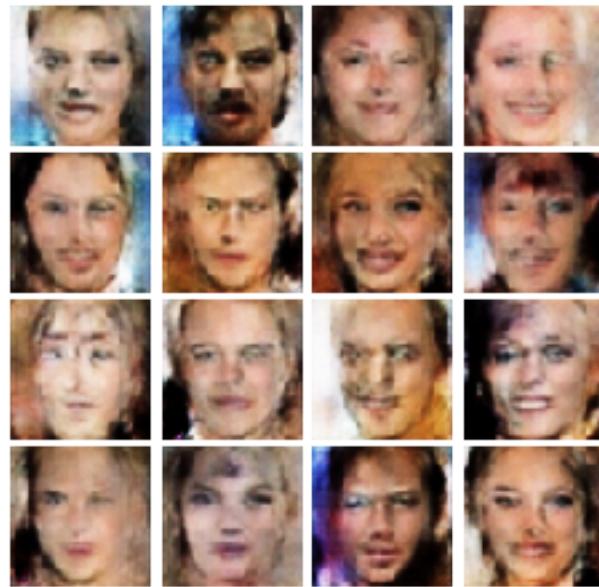
Iter: 3800, D: 0.1399, G: 0.3887



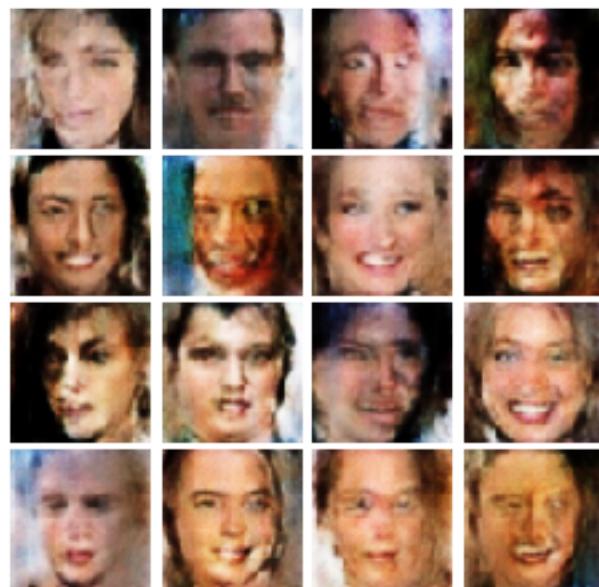
EPOCH: 5
Iter: 4000, D: 0.1381, G: 0.3465



Iter: 4200, D: 0.1666, G: 0.3524



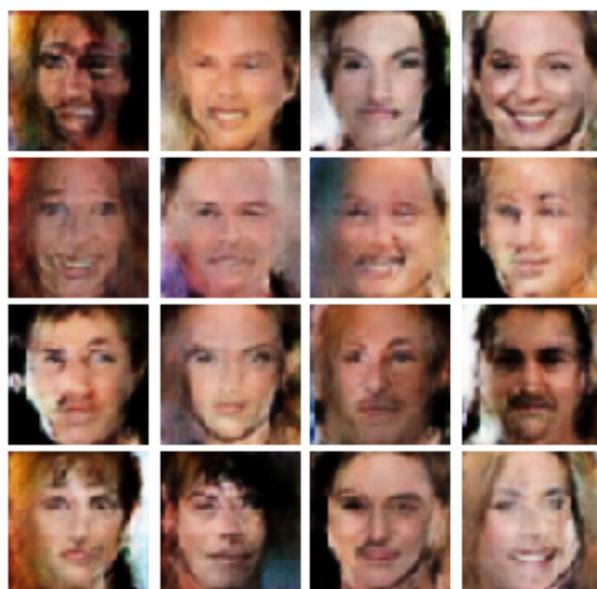
Iter: 4400, D: 0.1949, G:0.474



Iter: 4600, D: 0.4316, G:0.3245



Iter: 4800, D: 0.1164, G:0.7428



EPOCH: 6

Iter: 5000, D: 0.04333, G:0.4991



Iter: 5200, D: 0.1848, G: 0.2321



Iter: 5400, D: 0.3831, G: 0.2178



Iter: 5600, D: 0.1606, G:0.2322

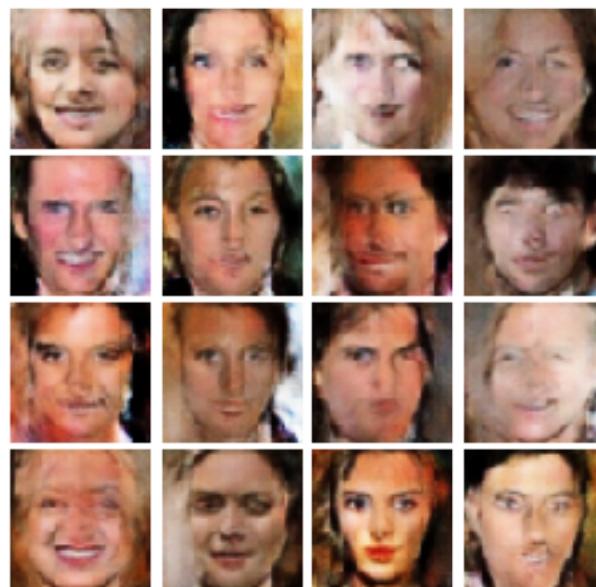


Iter: 5800, D: 0.09249, G:0.3506

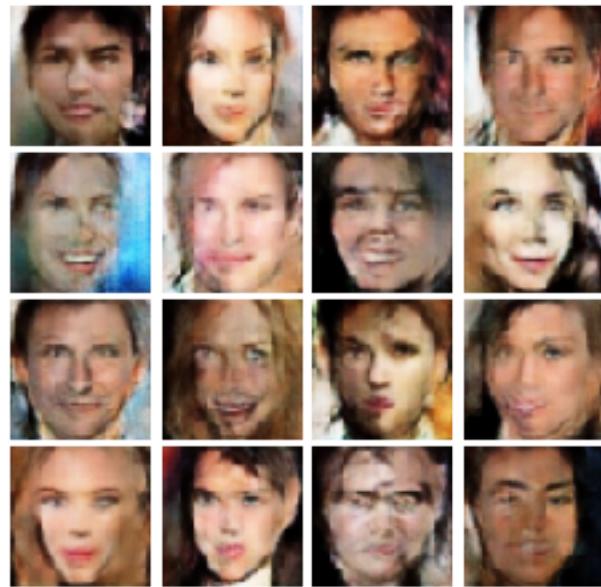


EPOCH: 7

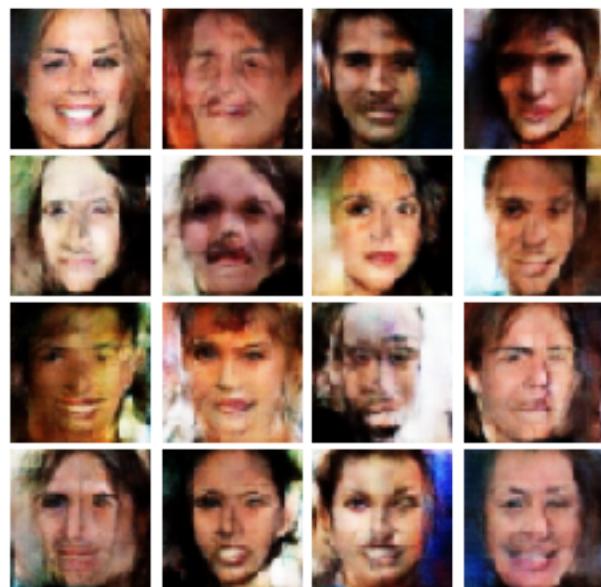
Iter: 6000, D: 0.1816, G:0.5065



Iter: 6200, D: 0.1331, G:0.2833



Iter: 6400, D: 0.2624, G:0.8327



Iter: 6600, D: 0.192, G:0.4704

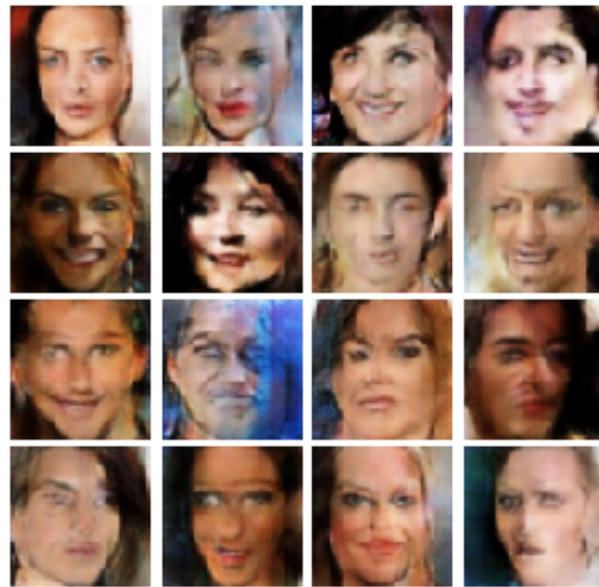


Iter: 6800, D: 0.0617, G: 0.4988

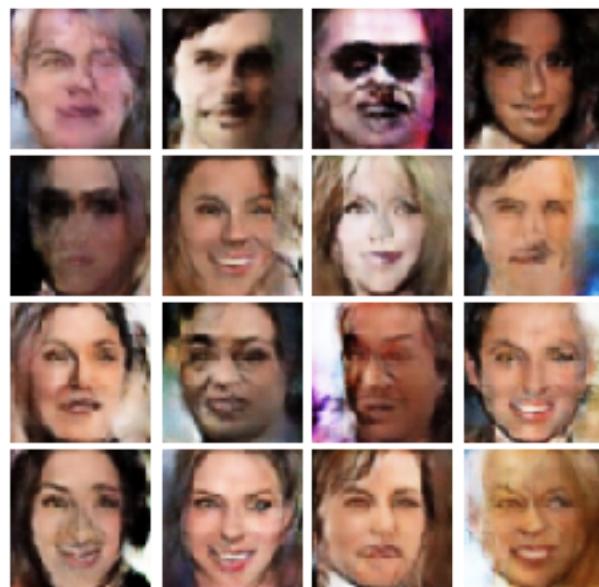


EPOCH: 8

Iter: 7000, D: 0.108, G: 0.3278



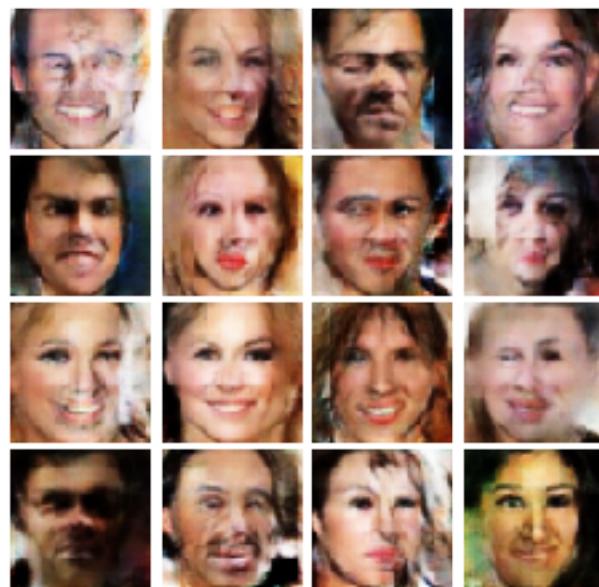
Iter: 7200, D: 0.116, G:0.3003



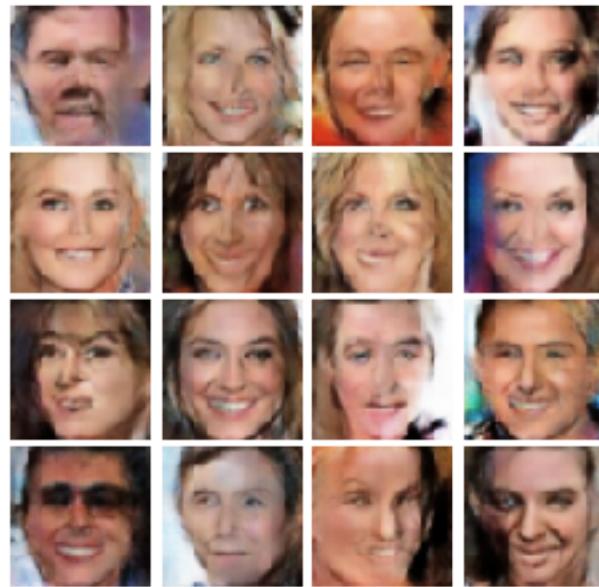
Iter: 7400, D: 0.1847, G:0.3113



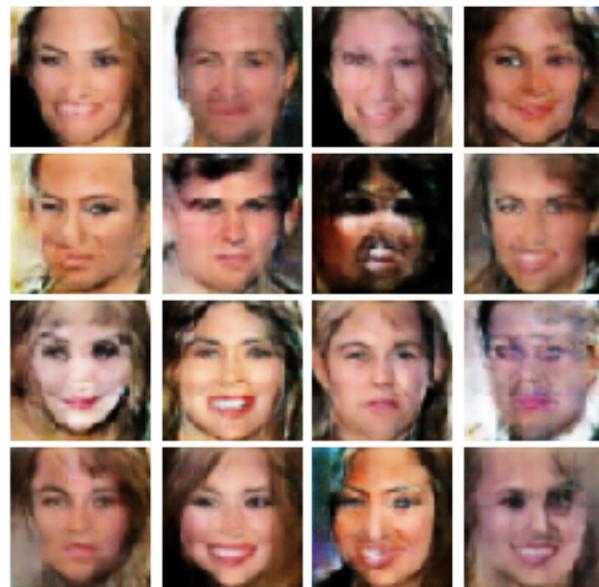
Iter: 7600, D: 0.1174, G:0.4725



Iter: 7800, D: 0.1766, G:0.5269



EPOCH: 9
Iter: 8000, D: 0.3225, G: 1.316



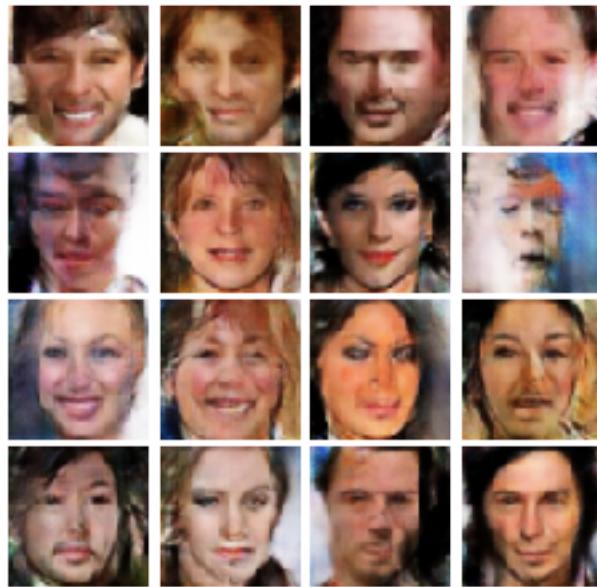
Iter: 8200, D: 0.06168, G: 0.3453



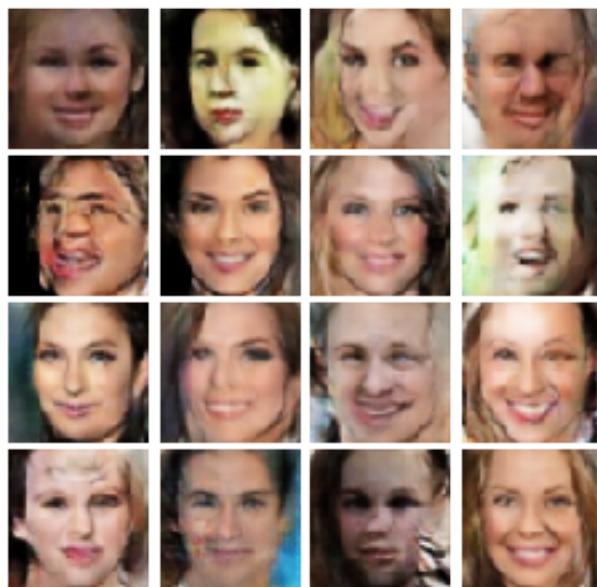
Iter: 8400, D: 0.07306, G: 0.5548



Iter: 8600, D: 0.2378, G: 0.3291



Iter: 8800, D: 0.08436, G:0.5598

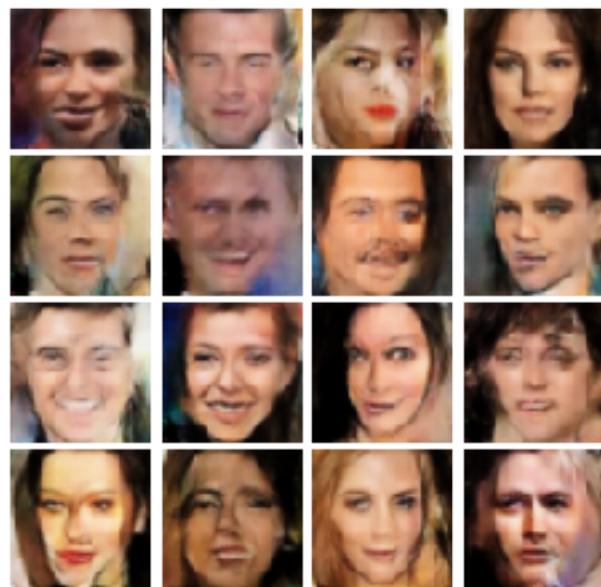


EPOCH: 10

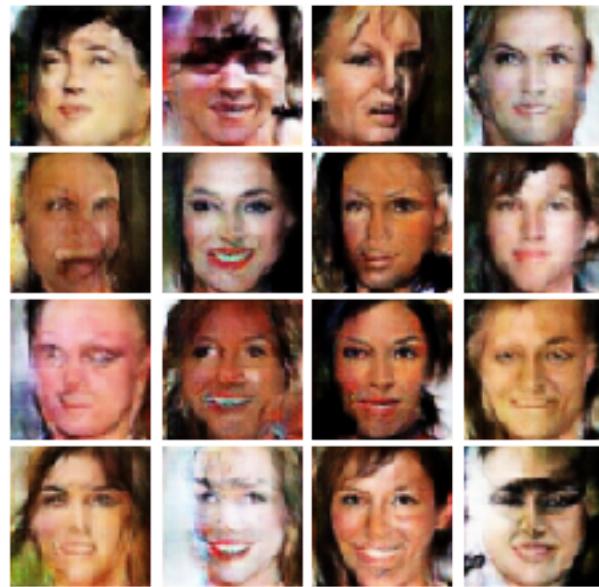
Iter: 9000, D: 0.06918, G:0.3911



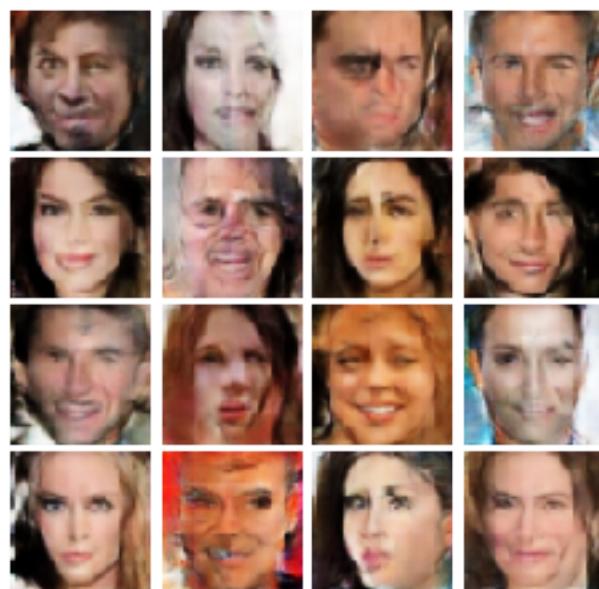
Iter: 9200, D: 0.05487, G:0.5438



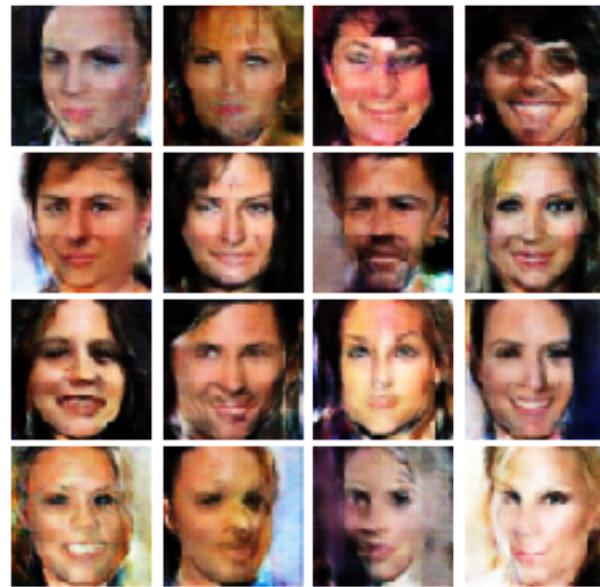
Iter: 9400, D: 0.09616, G:0.76



Iter: 9600, D: 0.1348, G:0.578

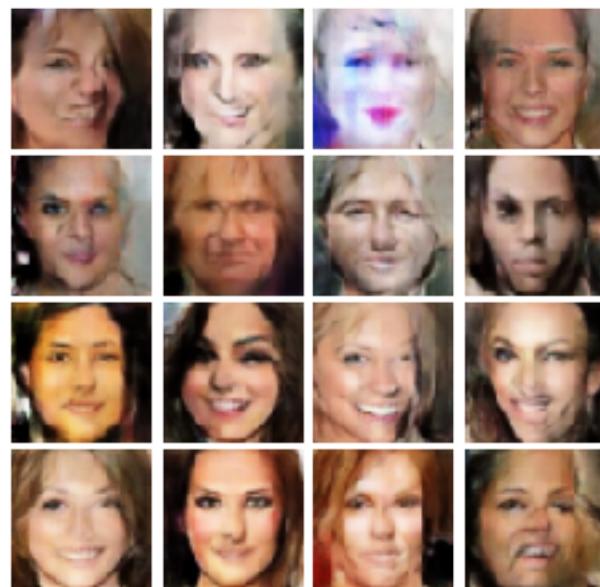


Iter: 9800, D: 0.05261, G:0.7094

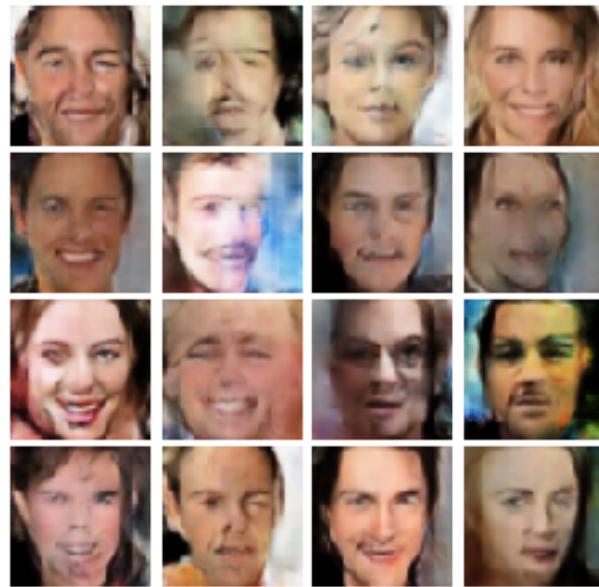


EPOCH: 11

Iter: 10000, D: 0.08958, G: 0.4279



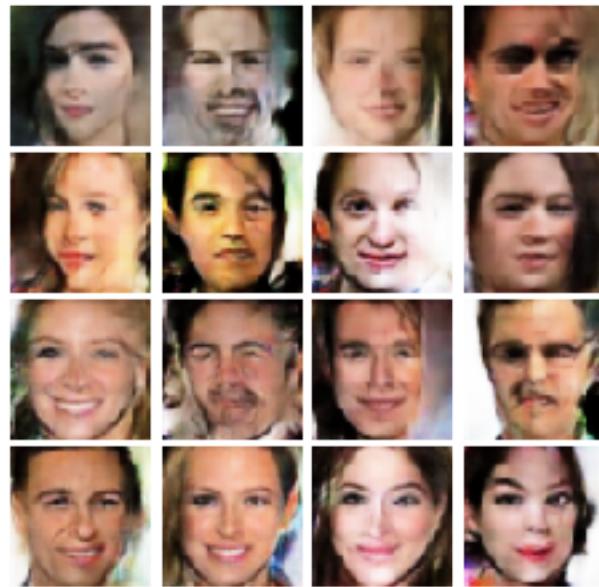
Iter: 10200, D: 0.3837, G: 0.09665



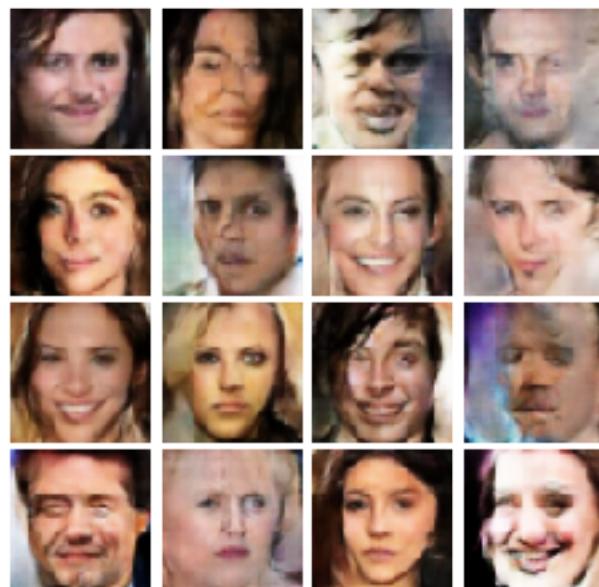
Iter: 10400, D: 0.08025, G: 0.4032



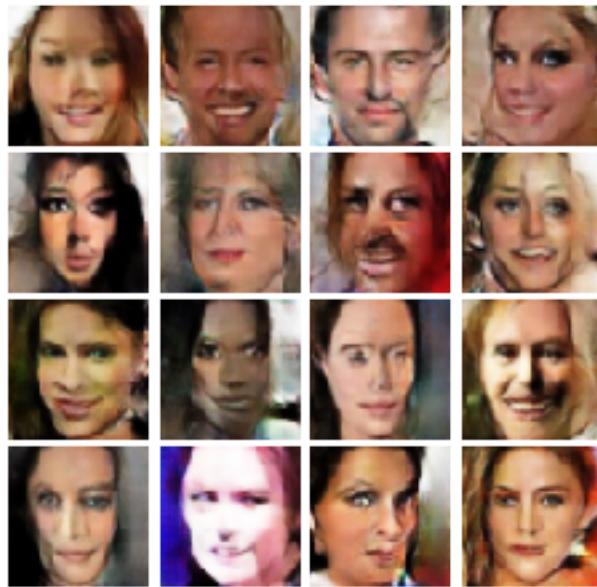
Iter: 10600, D: 0.156, G: 0.3459



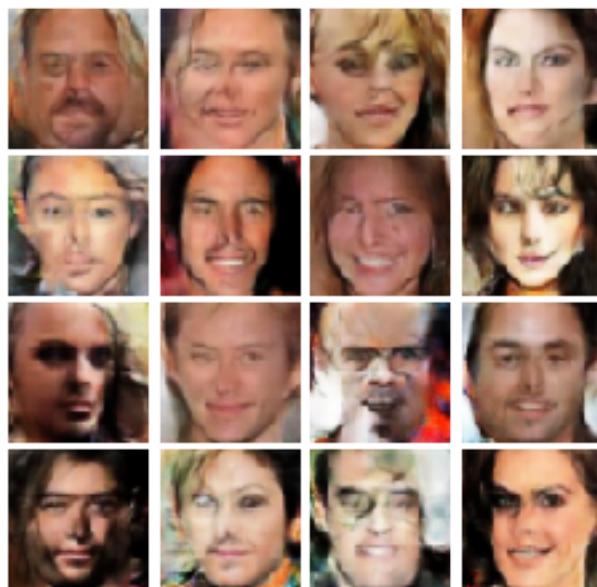
Iter: 10800, D: 0.1653, G: 0.389



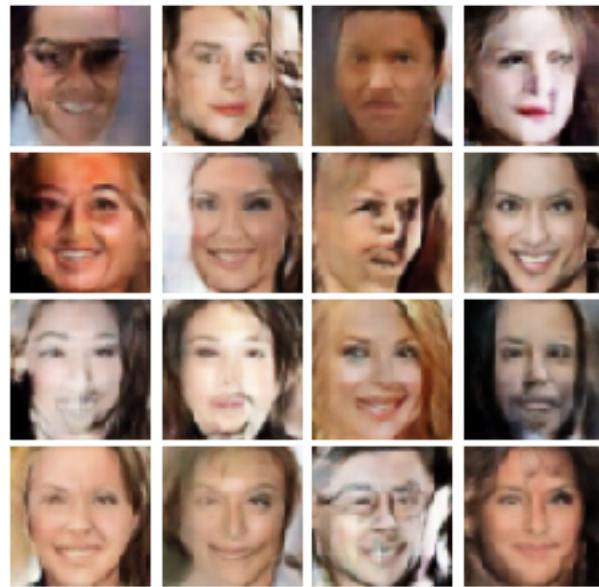
EPOCH: 12
Iter: 11000, D: 0.06717, G: 0.3078



Iter: 11200, D: 0.1218, G: 0.2829



Iter: 11400, D: 0.03085, G: 0.3041



Iter: 11600, D: 0.02614, G: 0.3798

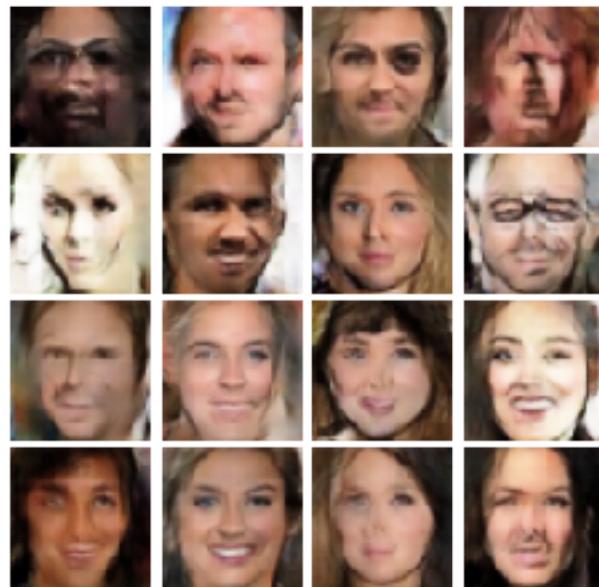


Iter: 11800, D: 0.2892, G: 0.2163



EPOCH: 13

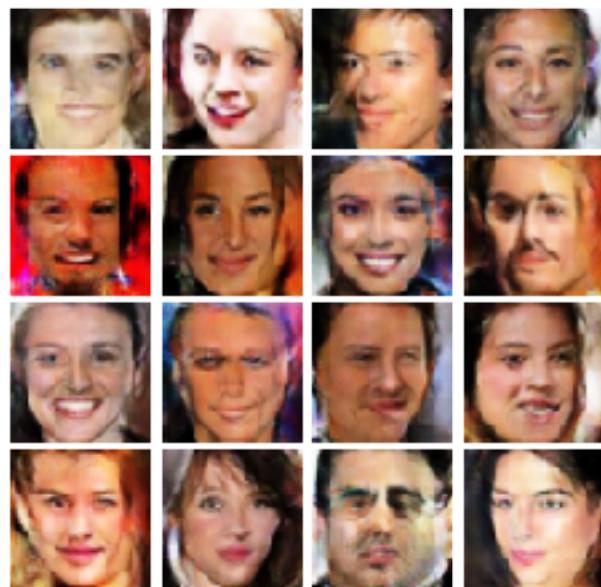
Iter: 12000, D: 0.1454, G: 0.3974



Iter: 12200, D: 0.1024, G: 0.4802



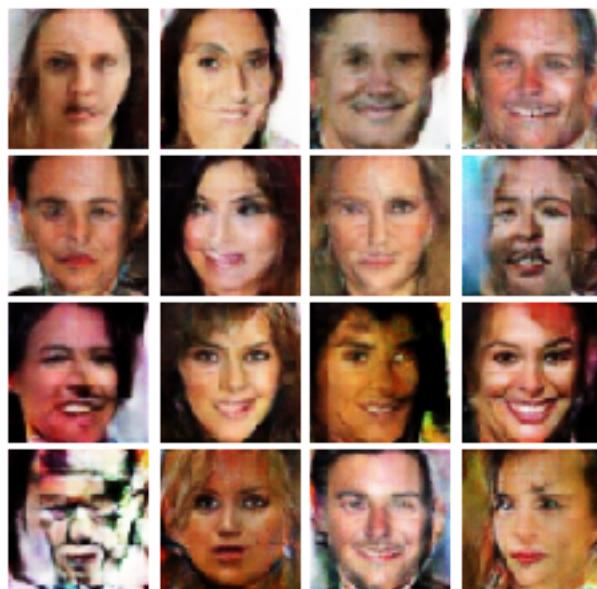
Iter: 12400, D: 0.07528, G: 0.4849



Iter: 12600, D: 0.0866, G: 0.3555



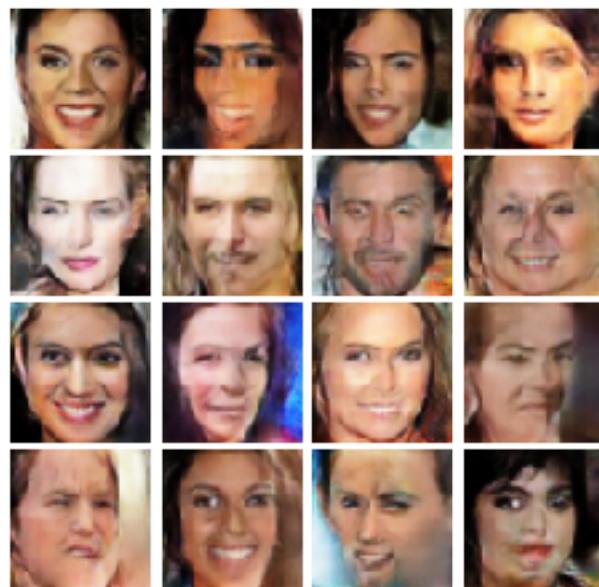
Iter: 12800, D: 0.1109, G: 0.3726



EPOCH: 14
Iter: 13000, D: 0.07298, G: 0.5951



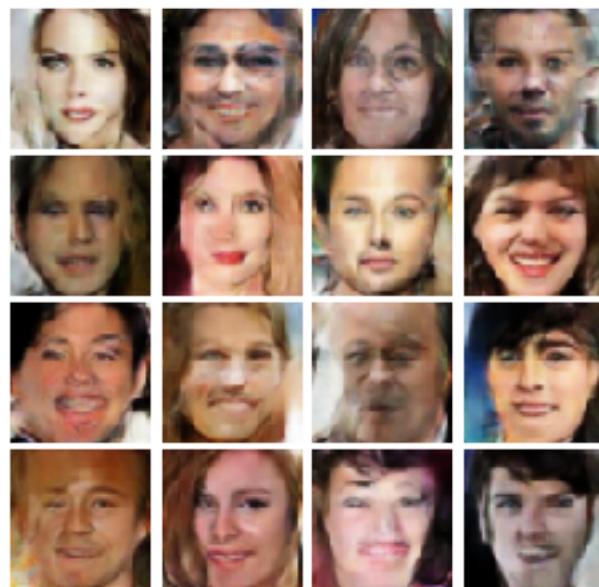
Iter: 13200, D: 0.08095, G: 0.2742



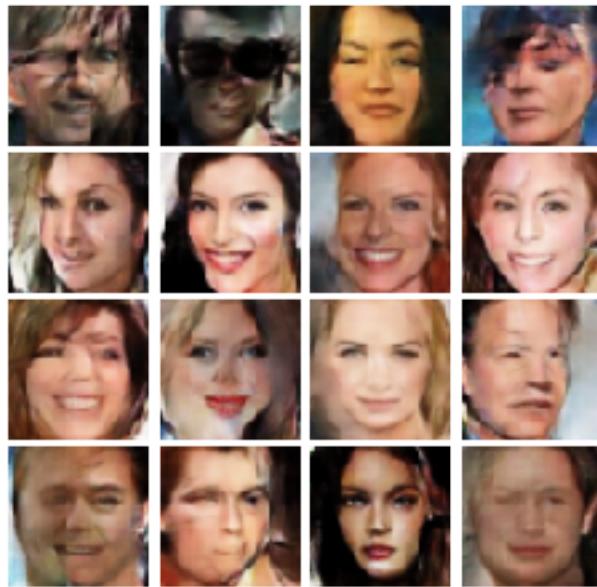
Iter: 13400, D: 0.05349, G: 0.2711



Iter: 13600, D: 0.07959, G:0.3187

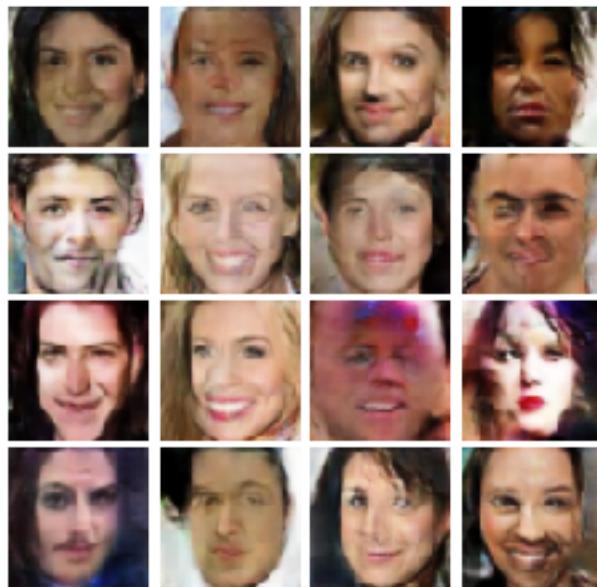


Iter: 13800, D: 0.09426, G:0.4375

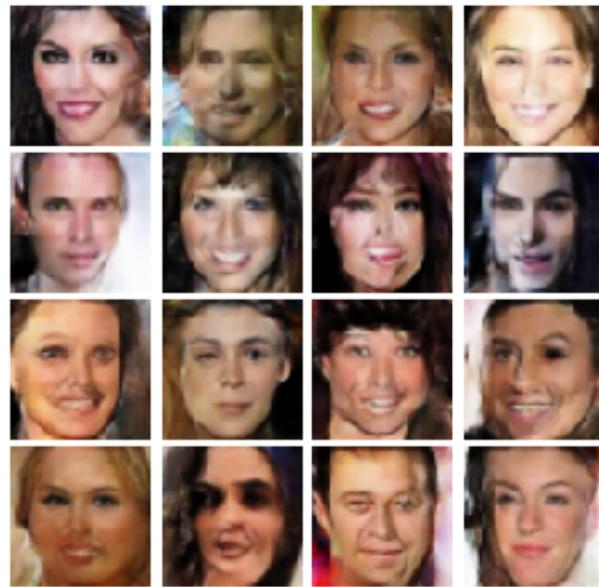


EPOCH: 15

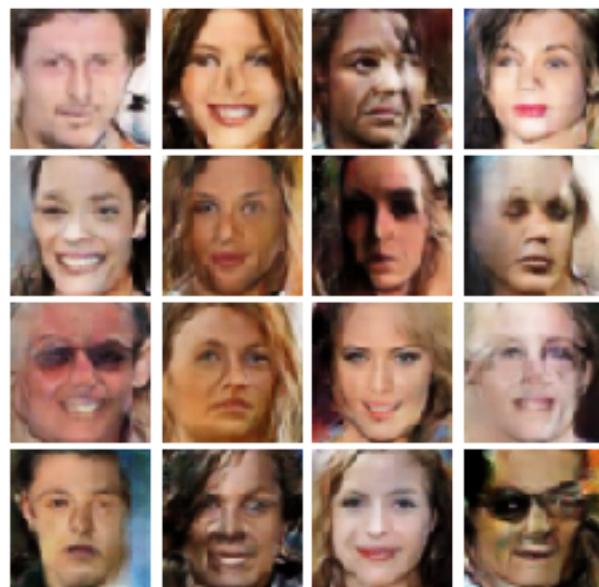
Iter: 14000, D: 0.07465, G: 0.6256



Iter: 14200, D: 0.1239, G: 0.5692



Iter: 14400, D: 0.04065, G: 0.2812



Iter: 14600, D: 0.03472, G: 0.341

